Trees and Ensemble Learning

FinTech

Lesson 11.2

Class Objectives

In today's class we'll learn about a new ML algorithms family: Tree based algorithms

Decision trees

Random forest

Weak learners

Ensemble methods

Categorical Data

Also we'll learn how to deal with categorical data

Color
Red
Red
Yellow
Green
Yellow

Red	Yellow	Green
1	0	0
1	0	0
0	1	0
0	0	1

Instructor Demonstration
Dealing with Text and Categorical Data in Machine Learning

Activity: Encoding Categorical Data for Machine Learning

In this activity, you will be tasked with encoding categorical and text features of a dataset that contains 2,097 loan applications.

Suggested Time:
10 minutes

Time's Up! Let's Review.

Walking into the Algorithms Forest

Tree based algorithms

Tree based algorithms, are part of the supervised machine learning methods.

Dimensionality Reduction

Meaningful Compression
Big Data Visualization
Structure Discovery
Feature Elicitation

Unsupervised Learning

Clustering

Recommender Systems
Targeted Marketing
Customer Segmentation

Machine Learning

Reinforcement Learning

Real-time Decisions
Robot Navigation
Game AI
Skill Acquisition

Classification

Image Classification
Identity Fraud Detection
Customer Retention
Diagnostics

Supervised Learning

Regression

Population Growth Prediction
Advertising Popularity Prediction
Weather Forecasting
Market Forecasting
Estimating Life Expectancy

Tree-based algorithms are supervised learning methods that are mostly used for classifications and regression problems.

# Tree Based Algorithms at a Glance

Bagging focuses on resampling data and running with different models on the fly in order to formulate the most accurate and precise prediction.

Boosting is both a process and set of algorithms. Boosting is the process of combining a set of weak learners into a stronger one.

Decision Trees

Bagging

Random Forest

Boosting

Gradient Boosting

A graphical representation of possible solutions to a decision based on certain conditions.

Bagging-based algorithm where only a subset of features are selected at random to build a forest or collection of decision trees.

Algorithm through parallel processing, tree-pruning, handling missing values and regularization to avoid overfitting/bias.

Algorithms

These algorithms can be used to solve classification or regression problems.
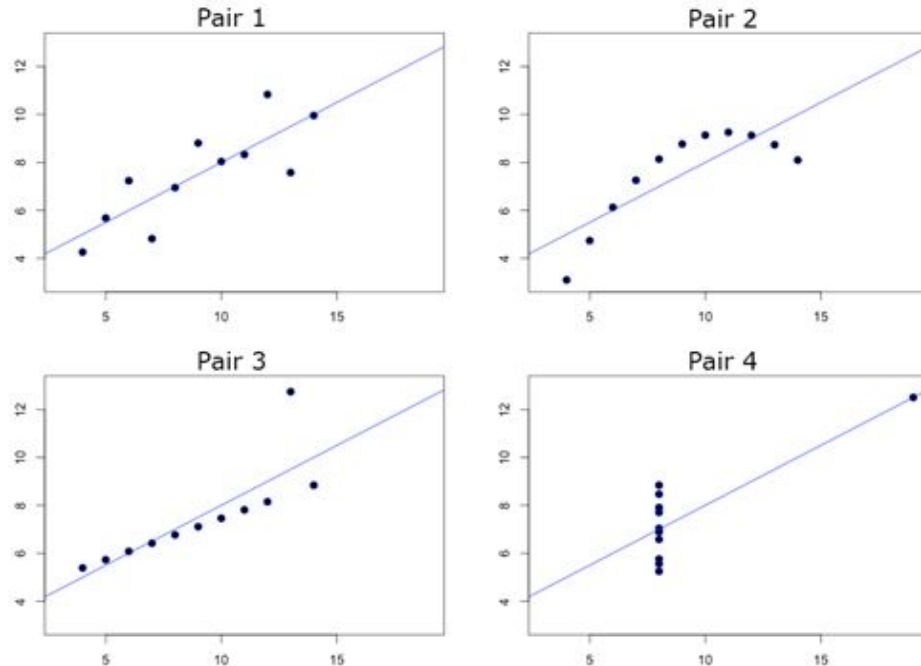
Classification

Regression

Linear vs. Non-Linear Models

In linear models, the relationship among input variables can be represented as a straight line, while non-linear models have a different shape.

Linear Models

Predicting the price of a house based on its size is an example of a linear problem. This is because, as a general rule, the size of the house is directly proportional to the price of the house.

Non-Linear Models

Tree-based algorithms can map non-linear relationships in data.

Non-Linear Models

Predicting if a credit application is going to be fraudulent or not may be an example of a non-linear problem due to the complex relationship between the input features and the output prediction.

Tree-Based Algorithms

These algorithms are quite often used in finance for assessing risk, preventing fraud, or fighting money laundering.

sklearn

`sklearn` has two modules that implement tree-based algorithms that we will be covering Today.
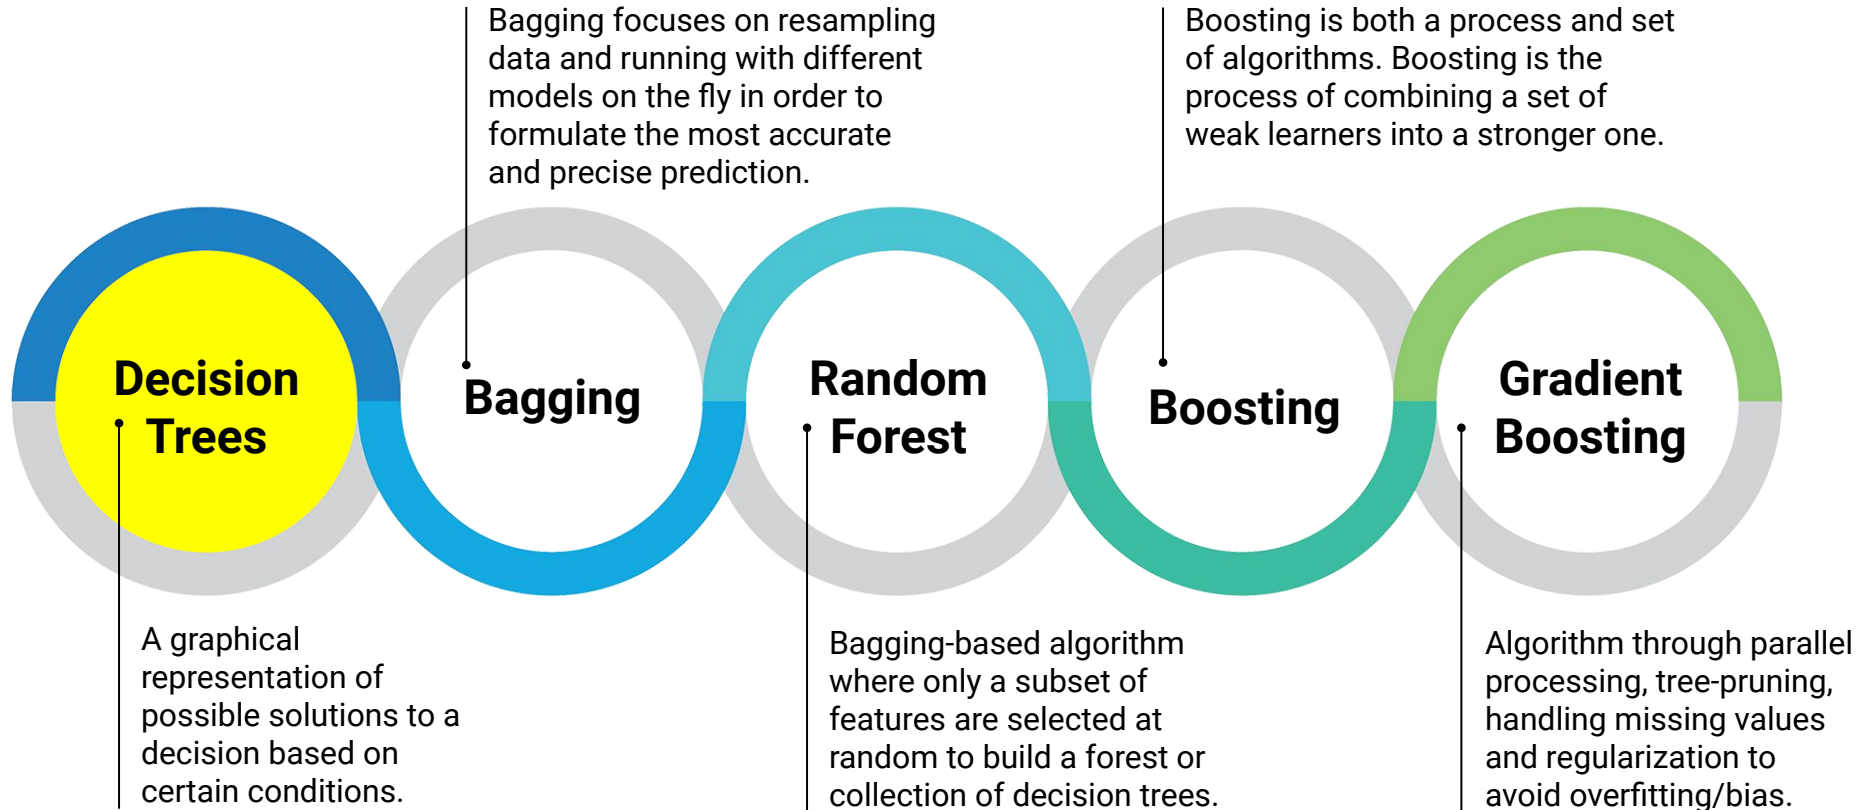
01	`sklearn.tree`

`sklearn.tree` implements decision trees.

| 02	`sklearn.ensemble`

`sklearn.ensemble` offers implementations for random forest, gradient boosting, boosting and bagging algorithms.

Decision Trees

Decision Trees

A graphical representation of possible solutions to a decision based on certain conditions.

Bagging

Bagging focuses on resampling data and running with different models on the fly in order to formulate the most accurate and precise prediction.

Random Forest

Bagging-based algorithm where only a subset of features are selected at random to build a forest or collection of decision trees.

Boosting

Boosting is both a process and set of algorithms. Boosting is the process of combining a set of weak learners into a stronger one.

Gradient Boosting

Algorithm through parallel processing, tree-pruning, handling missing values and regularization to avoid overfitting/bias.

Decision Trees

Decision trees encode a series of true/false questions.

Pet Picker:
Do you travel?

YES

NO

Are you gone more than one week per month?

Do you like to dress up your pets?

YES

NO

YES

NO

Decision Trees

These true/false questions can be represented with a series of if/else statements

Do you travel?

Yes Travel:

Are you gone for more than one week per month?

Yes: Pet Rock

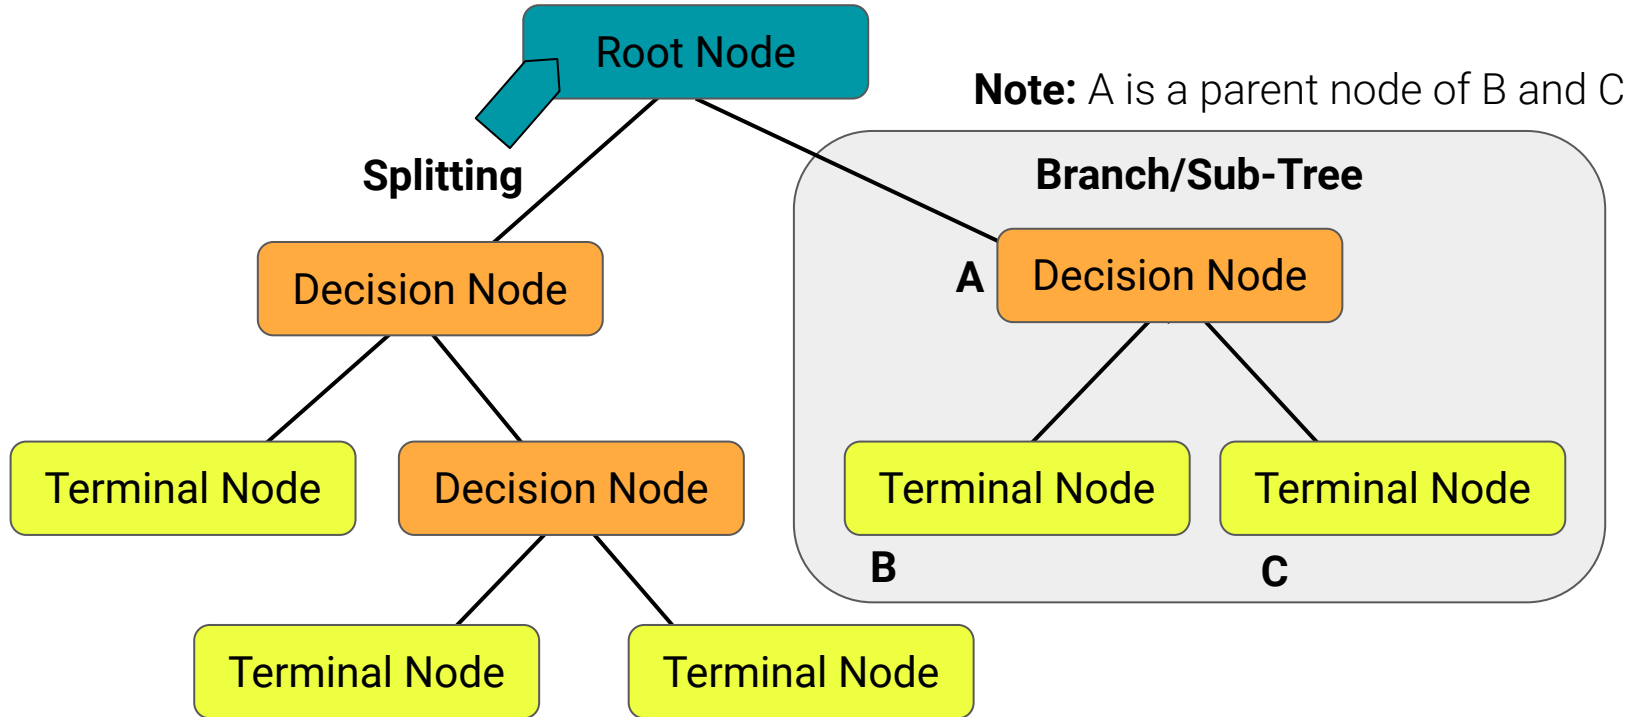No: Pet Fish

No Travel:

Do you like to dress up your pet?

Yes Dress Up: Pet Dog

No Dress Up: Pet Cat

```python
if (travel):
  if (time > week):
    print("Rock")
  else:
    print("Fish")
else:
  if (dress_up):
    print("Dog")
  else:
    print("Cat")
```

Root Node

Represents the entire population or sample data, this node gets divided into two or more homogeneous sets.
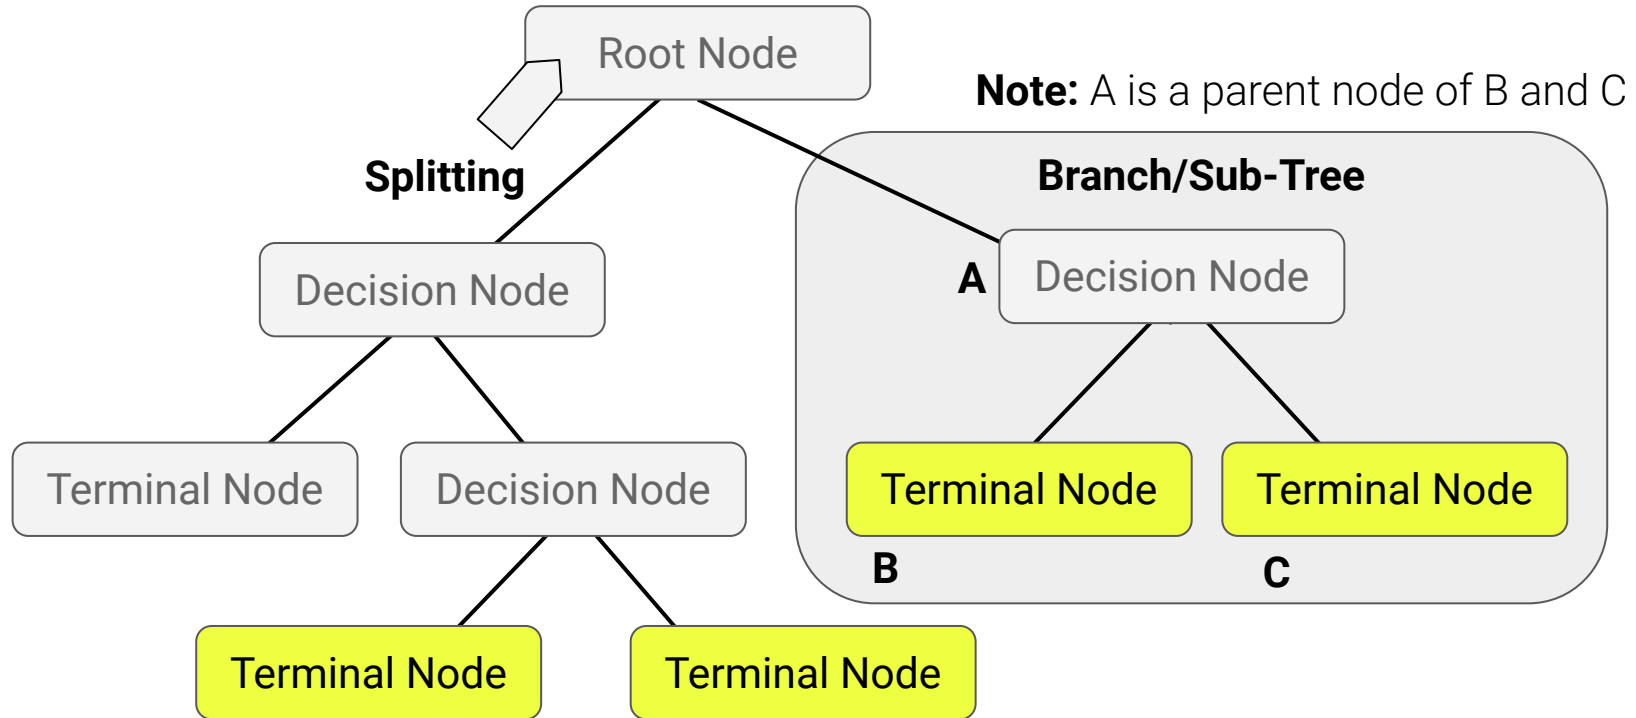
Parent Node

A node that is divided into sub-nodes.

Note: A is a parent node of B and C

Child Node

Sub-nodes of a parent node.

Root Node

Note: A is a parent node of B and C

Splitting

Branch/Sub-Tree

Decision Node

A Decision Node

Terminal Node

Decision Node

Terminal Node

B

Terminal Node

C

Terminal Node

Terminal Node

Decision Node

A sub-node that is split into further sub-nodes.

Root Node

Note: A is a parent node of B and C

Splitting

Branch/Sub-Tree

Decision Node

A Decision Node

Terminal Node

Decision Node

Terminal Node

Terminal Node

B

C

Terminal Node

Terminal Node

Decision Trees

Key concepts to know while working with decision trees:

Leaf or Terminal Node	Nodes that do not split.
Branch or Sub-Tree	A subsection of entire tree.
Splitting	Process of dividing a node into two or more sub-nodes.
Pruning	Process of removing sub-nodes of a decision node.
Tree's Depth	The number of decision nodes encountered before making a decision.

Decision Trees

Decision trees can become very complex and may not generalize well.
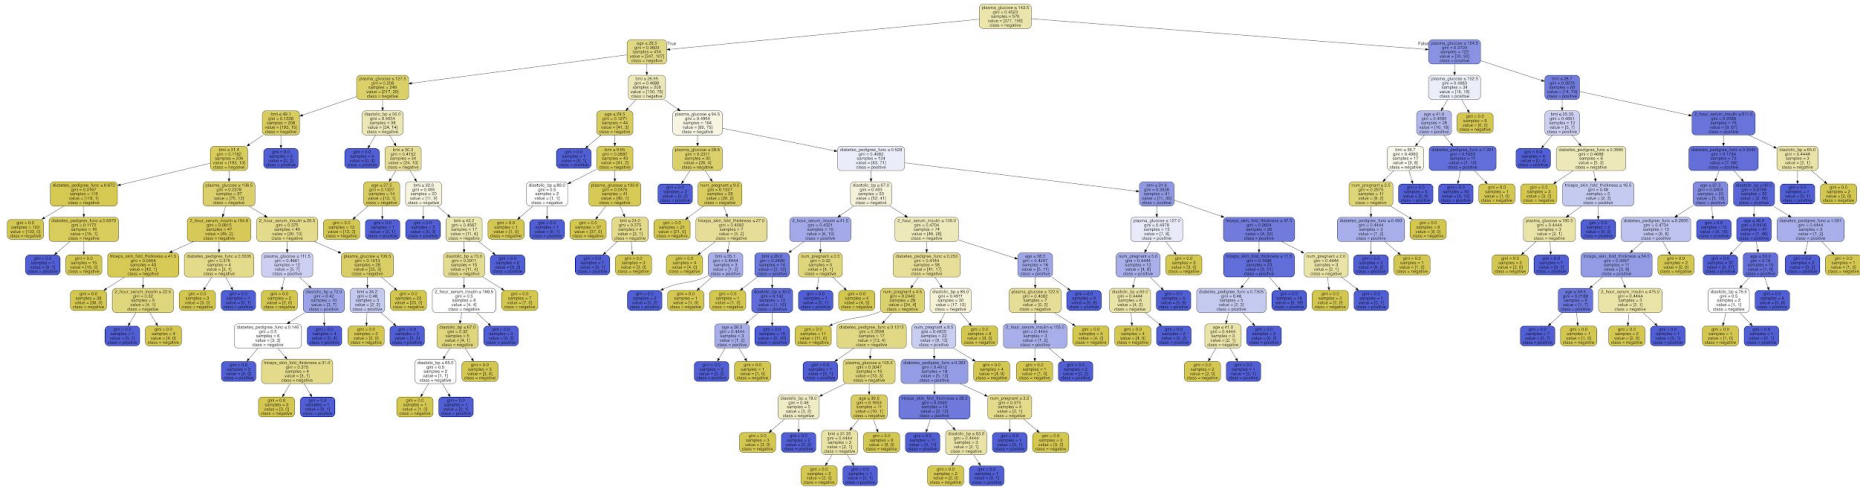
Instructor Demonstration
Decision Trees

Activity: Predicting Fraudulent Loans Applications

In this activity, you will create a decision tree model to predict fraudulent loan applications.

Time's Up! Let's Review.

Introduction to Ensemble Learning

The Classification Algorithm Race

If we compare the performance of classification algorithms, we'll find that some algorithms performed better than others

Weak Learners

- Algorithms that actually fail at learning in an adequate fashion.

- They are a consequence of limited data to learn from.

- Their predictions are only a little better than random chance.

Weak Learners are still valuable in Machine Learning

They can be combined with other classifiers in order to make a more accurate and robust prediction engine.

Combined weak learners are an example of ensemble learning:

Ensemble Learners

Ensemble learners improves accuracy and robustness, as well as decrease variance.

Combined, weak learners can perform as well as strong learners.

Combining Weak Learners

Weak learners have to be combined using specific algorithms like:
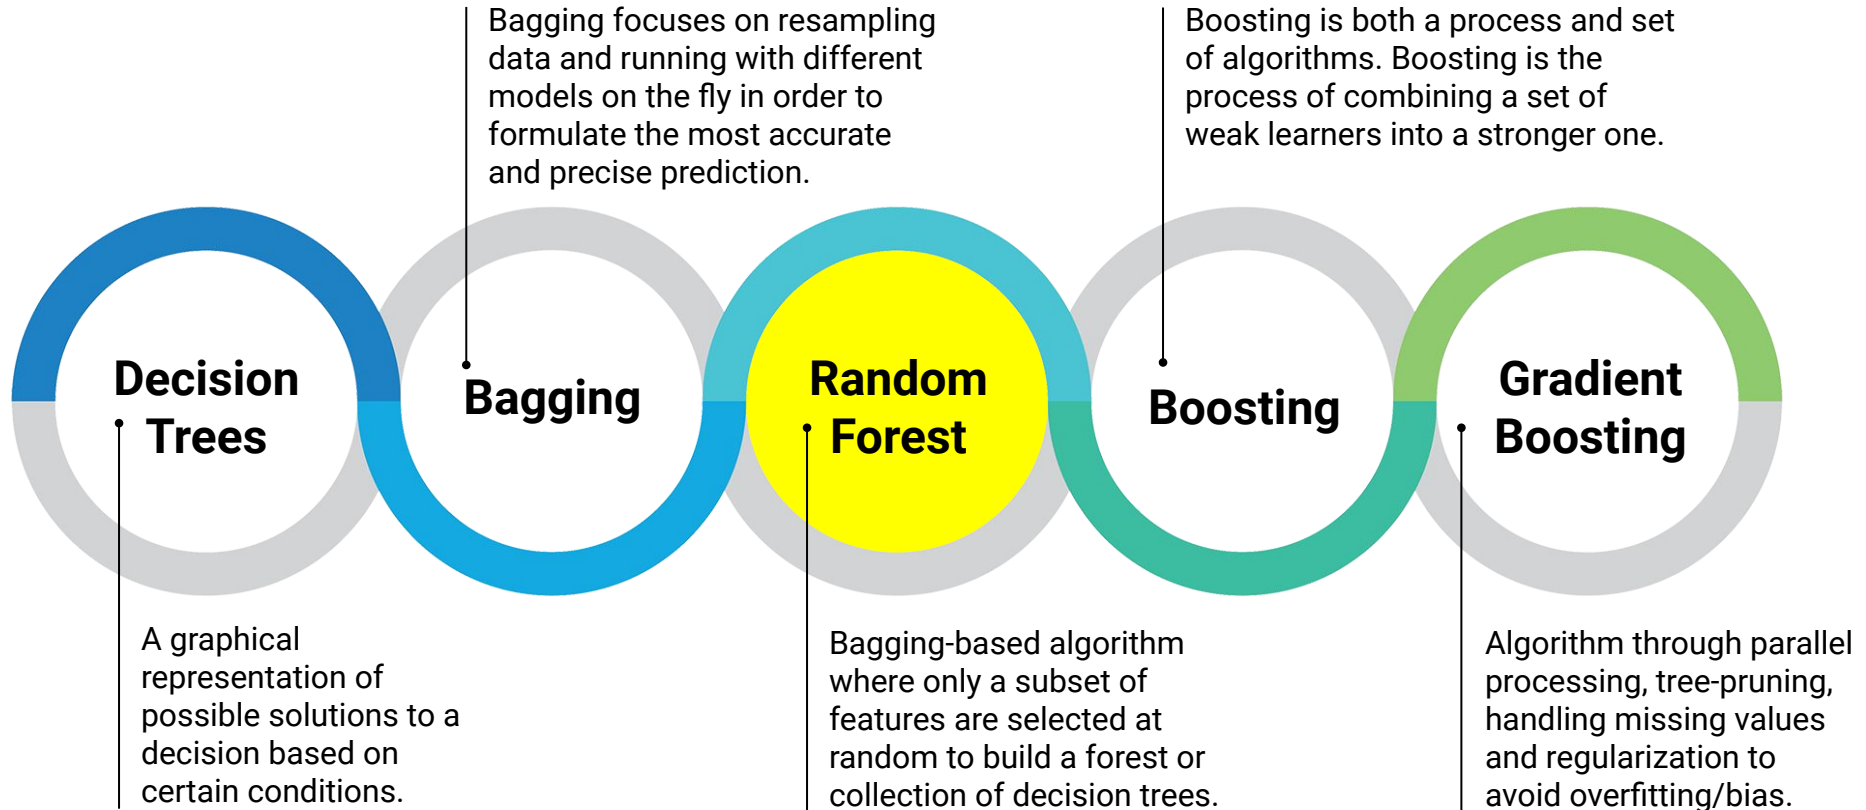
GradientBoostingTree

XGBoost

Random Forest

Random Forest

Bagging focuses on resampling data and running with different models on the fly in order to formulate the most accurate and precise prediction.

Boosting is both a process and set of algorithms. Boosting is the process of combining a set of weak learners into a stronger one.

Decision Trees

Bagging

Random Forest

Boosting

Gradient Boosting

A graphical representation of possible solutions to a decision based on certain conditions.

Bagging-based algorithm where only a subset of features are selected at random to build a forest or collection of decision trees.

Algorithm through parallel processing, tree-pruning, handling missing values and regularization to avoid overfitting/bias.

Random Forest

Instead of having single, complex tree like the ones created by decision trees, a random forest algorithm will sample the data and build several smaller, simpler decisions trees.
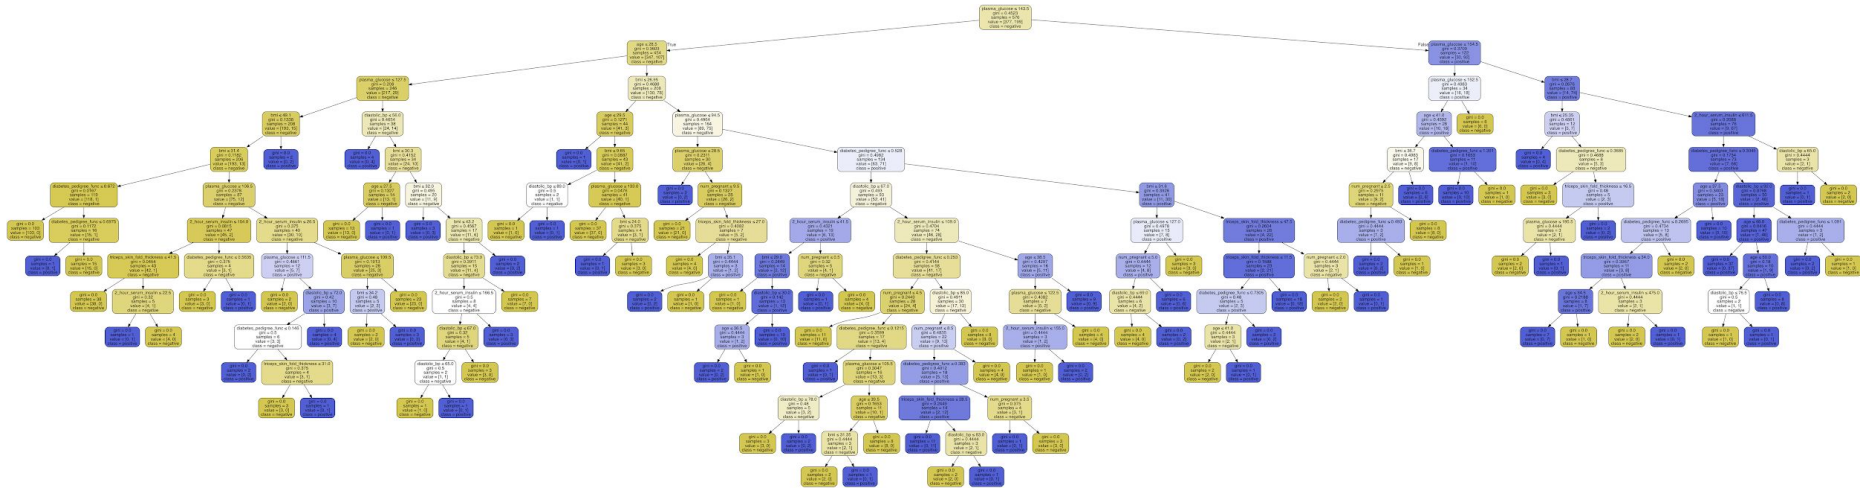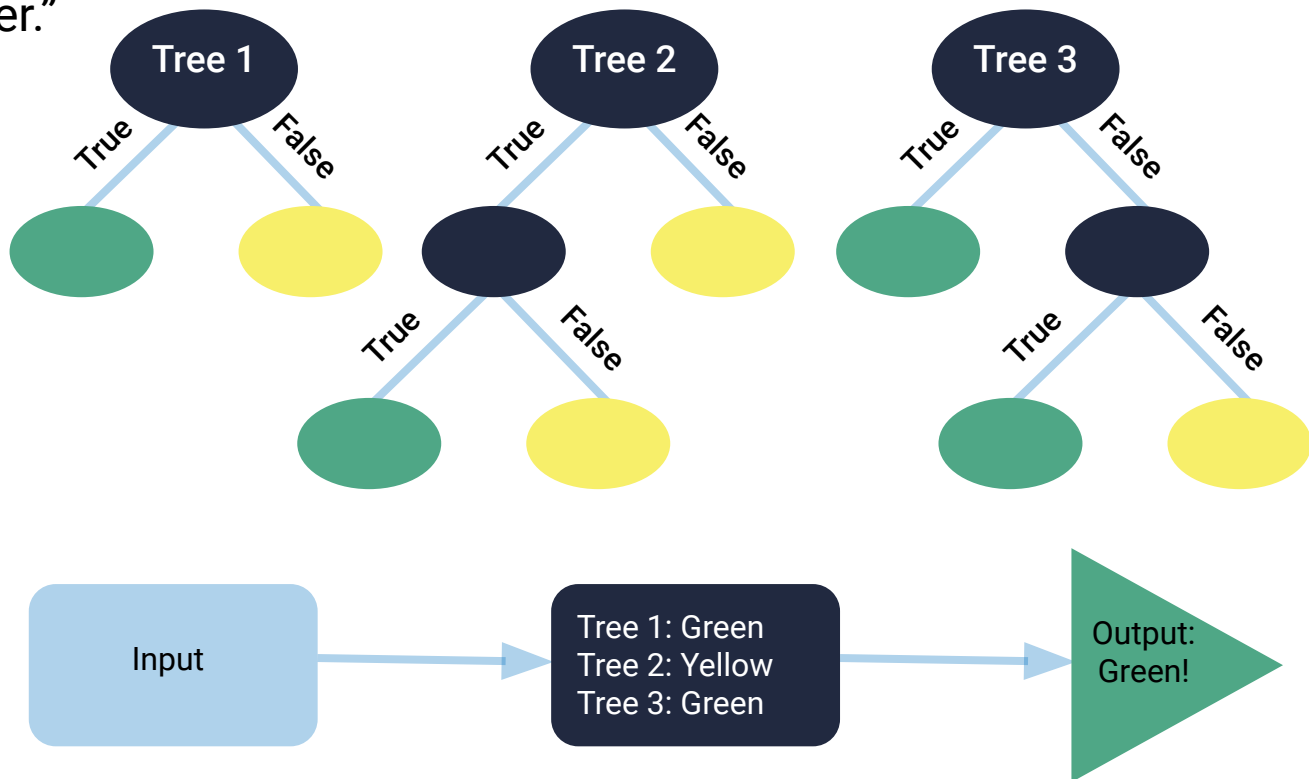
A Forest of Trees

In a random forest, each tree is much simpler because it is built from a subset of the data.

A Forest of Trees

Each tree is considered a "weak classifier" but when you combine them, they form a "strong classifier."

Benefits of Random Forest Algorithm

It's robust against overfitting.

It can be used to rank the importance of input variables in a natural way.

It can handle thousands of input variables without variable deletion.

It's robust to outliers and non-linear data.

It runs efficiently on large databases.

Instructor Demonstration
Random Forest

Activity: Predicting Fraud with Random Forests

In this activity, you will explore how the random forest algorithm can be used to identify fraudulent loan applications. You will use the `sba_loans_encoded.csv` file that they created before to train the model.

Suggested Time:
10 minutes

Time's Up! Let's Review.

Review: Predicting Fraud with Random Forests

Would you trust in this model to deploy a fraud detection solution in a bank?

What are your insights about the top 10 most importance features?

Boosting and Bagging

Boosting and Bagging

Bagging focuses on resampling data and running with different models on the fly in order to formulate the most accurate and precise prediction.

Boosting is both a process and set of algorithms. Boosting is the process of combining a set of weak learners into a stronger one.

Decision Trees

Bagging

Random Forest

Boosting

Gradient Boosting

A graphical representation of possible solutions to a decision based on certain conditions.

Bagging-based algorithm where only a subset of features are selected at random to build a forest or collection of decision trees.

Algorithm through parallel processing, tree-pruning, handling missing values and regularization to avoid overfitting/bias.

Boosting and Bagging

Boosting and bagging algorithms are used to improve the robustness and reliability of machine-learning models


Bagging and Boosting are both **ensemble methods** in Machine Learning

Bagging

Random sampling with replacement

Boosting

Random sampling with replacement over weighted data

Boosting and Bagging

Boosting and bagging algorithms like XGBoost are often the best performing in Kaggle machine-learning contests. Their ability to make accurate predictions with precision and substantial recall is almost unparalleled

12 Competitions

Santander Value Prediction Challenge
Predict the value of transactions for potential customers.
Featured · a year ago · 🏷 finance, banking

$60,000
4,477 teams

Two Sigma Financial Modeling Challenge
Can you uncover predictive value in an uncertain world?
Featured · Code Competition · 3 years ago · 🏷 finance, future prediction

$100,000
2,070 teams

The Winton Stock Market Challenge
Join a multi-disciplinary team of research scientists
Featured · 4 years ago · 🏷 finance, tabular data, future prediction

$50,000
832 teams

Boosting vs. Bagging

Boosting

Boosting takes multiple algorithms and coordinates them as an ensemble and runs the algorithms in tandem to identify the best prediction

VS.

Bagging

Bagging focuses on resampling data and running with different models on the fly to formulate the most accurate and precise prediction.

Bagging

Bagging averages predictions from multiple samples and or models.

Boosting

Boosting algorithms work iteratively, by sampling more heavily the observations with worst predictions. Subsequent weak learners then learn these more. Weak learners are then aggregated to produce a more accurate and precise prediction. The goal goal of a boosting algorithm is to combine weak learners into ensemble learners.
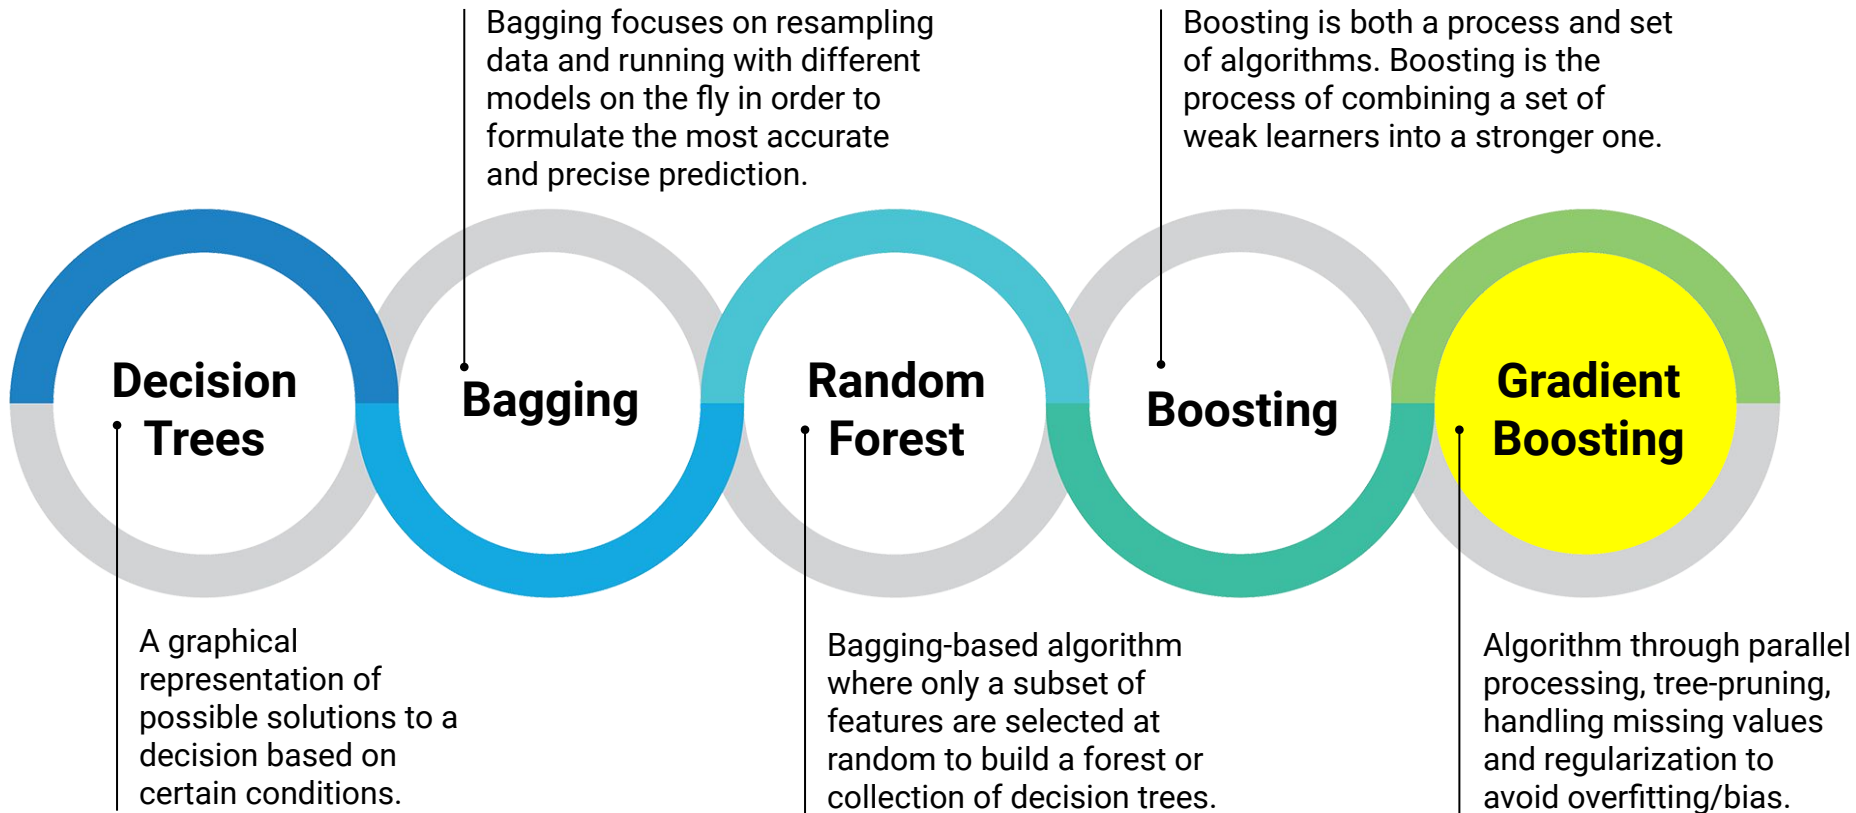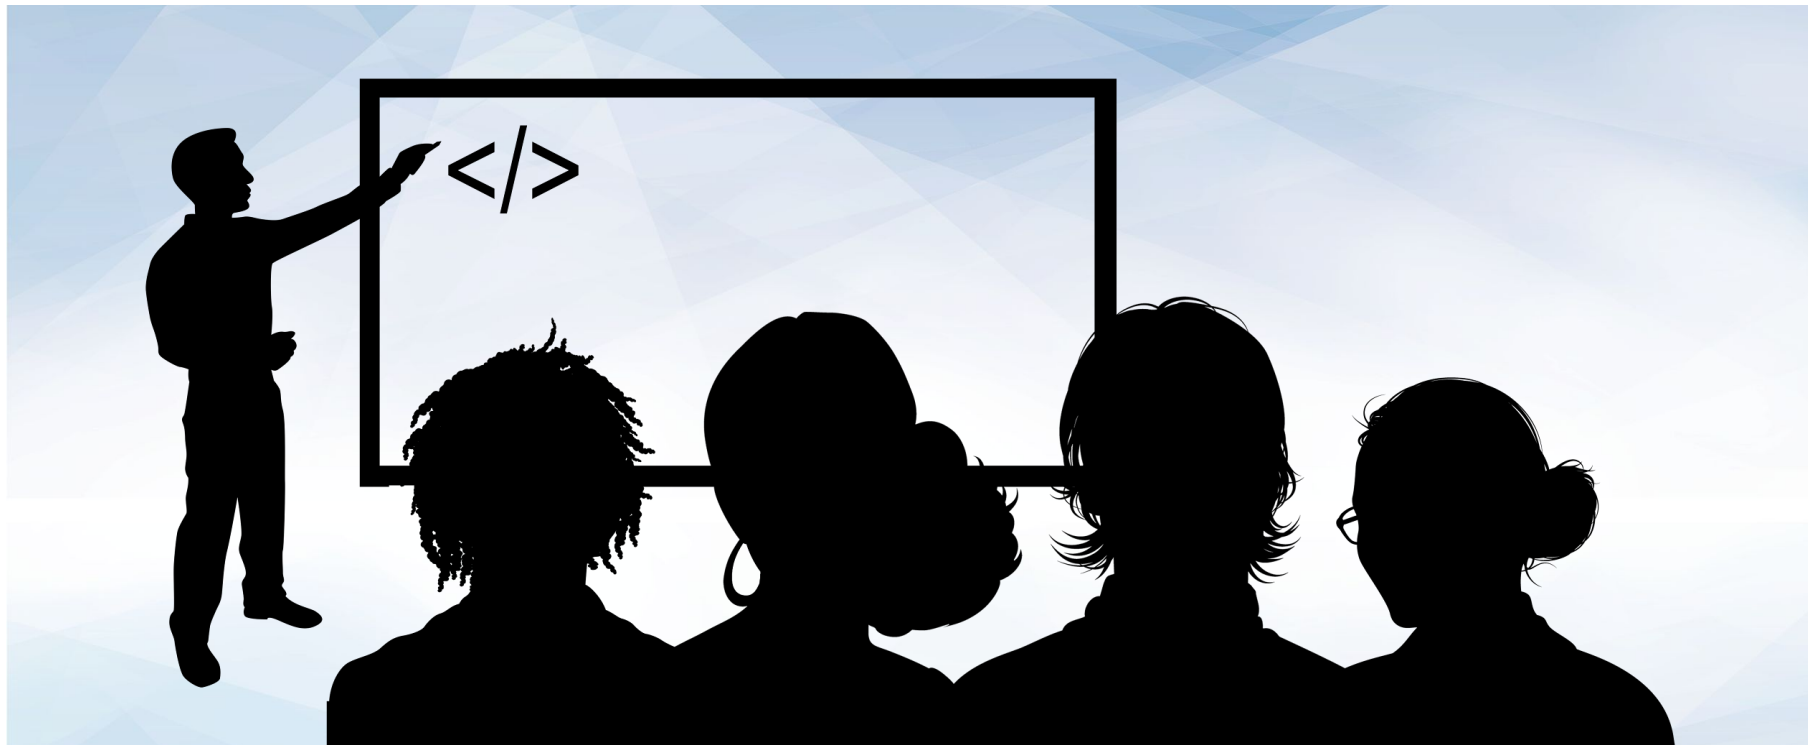
Boosting

For this reason, boosting algorithms are considered meta-algorithms. Instead of working with and affecting data, boosting algorithms work with and affect other algorithms.

Gradient Boosted Tree

Gradient Boosting

Bagging focuses on resampling data and running with different models on the fly in order to formulate the most accurate and precise prediction.

Boosting is both a process and set of algorithms. Boosting is the process of combining a set of weak learners into a stronger one.

Decision Trees

Bagging

Random Forest

Boosting

Gradient Boosting

A graphical representation of possible solutions to a decision based on certain conditions.

Bagging-based algorithm where only a subset of features are selected at random to build a forest or collection of decision trees.

Algorithm through parallel processing, tree-pruning, handling missing values and regularization to avoid overfitting/bias.

Instructor Demonstration
Gradient Boosted Tree

Gradient Boosted Tree

Gradient Boosted Trees can be created using the GradientBoostingClassifier module from the ensemble package

From `sklearn.ensemble` import `GradientBoostingClassifier`

Gradient Boosted Tree

Arguments

GradientBoostingClassifier has three main arguments:

`N_estimators`

`Learning_rate`

`Max_depth`

Definitions

The `n_estimators` parameter configures the number of weak learners being used with the boosting algorithm.

Gradient Boosted Tree

`Learning_rate`

`Learning_rate` controls overfitting.

Smaller values should be used when setting learning_rate.

`max_depth`

The `max_depth` argument identifies the size/depth of each decision tree being used.

Activity: Turbo Boost

In this activity you will use the sklearn `GradientBoostingClassifier` boosting algorithm to detect fraudulent loan applications using ensemble learning.

Time's Up! Let's Review.

The Trees vs. The World

Classification: A multidisciplinary challenge

Finance and Banking	Fraud detection, money laundering, credit risk assessment.
Retail and Marketing	Customized product offers, product recommendation, direct-marketing optimization.
Politics	Vote intention, party affinity.
Health	Trials tests, ills diagnosis.
Security	Intruders detection, predictive maintenance.
Education	Programs affinity, customized curricula, desertion prevention.

Classification: A multidisciplinary challenge
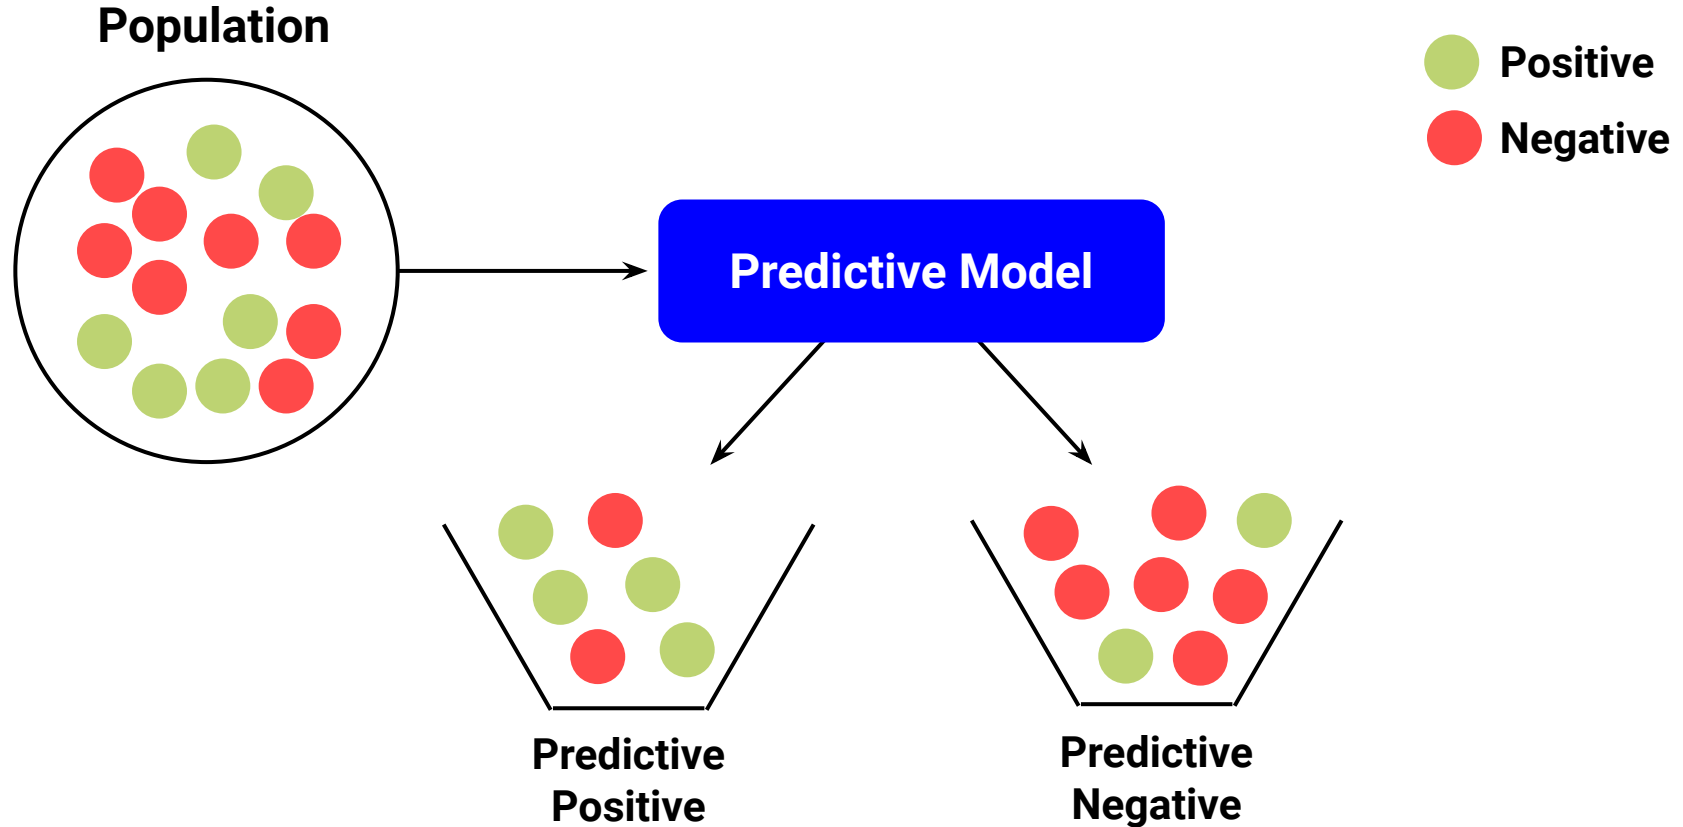
Are tree-based algorithms the strongest for classification?

Tree-based algorithms

Are easy to represent, making a complex model much easier to interpret.

Can be used for any type of data: Numerical (e.g., loan's amount) or categorical (e.g., name of bank that issues a loan).

Require little data preparation.

Can handle data that are not normally distributed.

Can avoid overfitting.

Trees vs. Classical Classifiers

Generally speaking, classical classifiers may be faster.

Logistic regression may outperform decision trees or random forests having a large number of features with low noise.

SVM also support linear and non-linear models.

SVM handles outliers better.

KNN naturally supports incremental learning (data streams).

Which algorithm should I use for classification?

Questions?