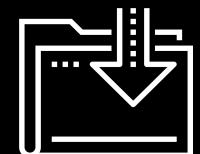


Recurrent Neural Networks

FinTech
Lesson 14.3v1.3





What Are We Going to Learn Today?

What We'll Learn Today

In this lesson, we'll cover the following:

01

RNN, the neural network that remembers the past

02

How to assess and compare models' performance

03

How to score sentiment using RNNs

04

How to use RNNs for time-series analysis



Introducing the ROC Curve and AUC

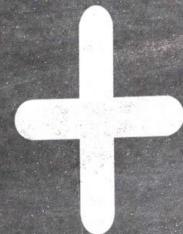
Introducing the ROC Curve and AUC

Recall that we use the confusion matrix to assess the performance of a binary classification model.

		Predicted Values	
		Positive (1)	Negative (0)
Actual Values	Positive (1)	TP	FN
	Negative (0)	FP	TN

Introducing the ROC Curve and AUC

Also recall the four components of this matrix:



TP

(True Positives)

Refers to the positive values that were correctly classified as positive by the model.

TN

(True Negatives)

Refers to the negative values that were correctly classified as negative by the model.

FP

(False Positives)

Refers to the negative values that were incorrectly classified as positive by the model.

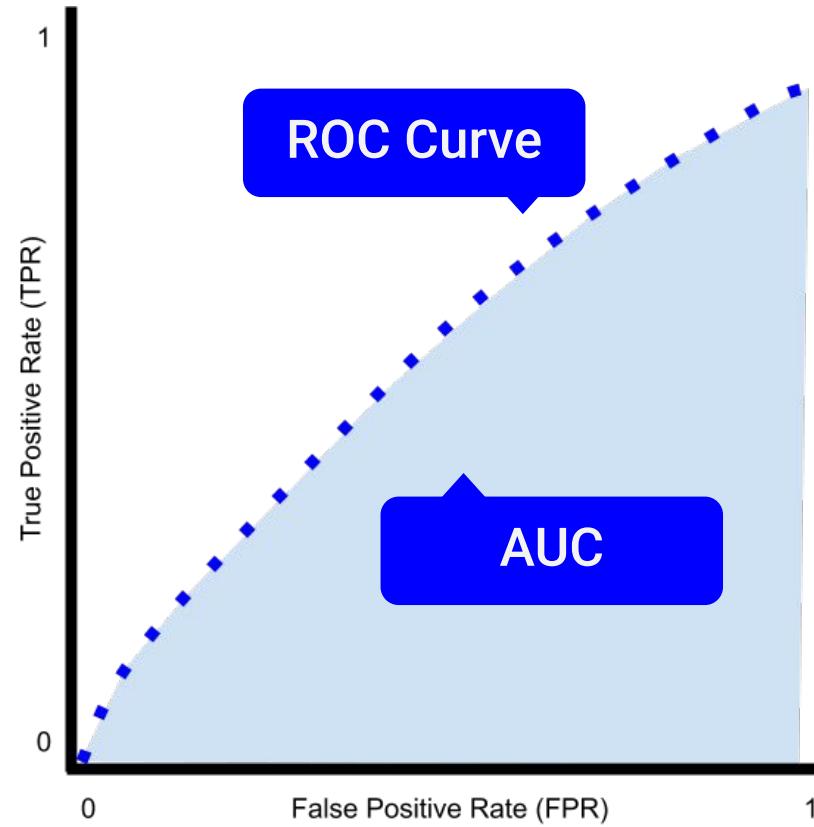
FN

(False Negatives)

Refers to the positive values that were incorrectly classified as negative by the model.

Introducing the ROC Curve and AUC

The ROC curve and AUC are techniques that allow us to check and visualize the performance of a classification model.



Understanding ROC

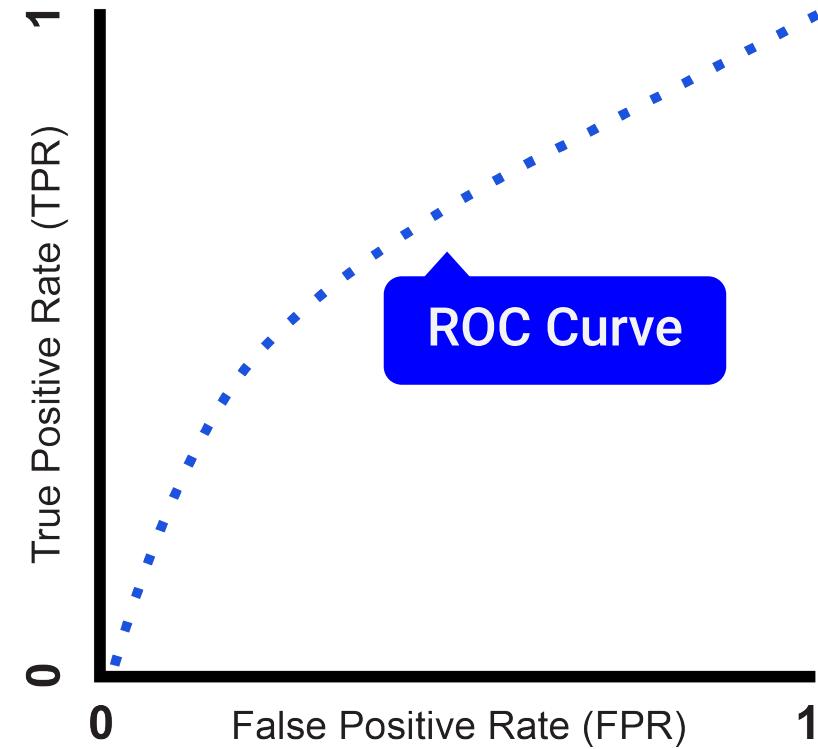
ROC

stands for

Receiver Operating Characteristic

Introducing the ROC Curve and AUC

The ROC curve shows the performance of a classification model as its discrimination threshold is varied.



Introducing the ROC Curve and AUC

To plot a ROC curve, we use two parameters:

TPR

True Positive Rate

(also known as recall)



The **TPR** is calculated as follows:

$$TPR = \frac{TP}{TP + FN}$$

FPR

False Positive Rate



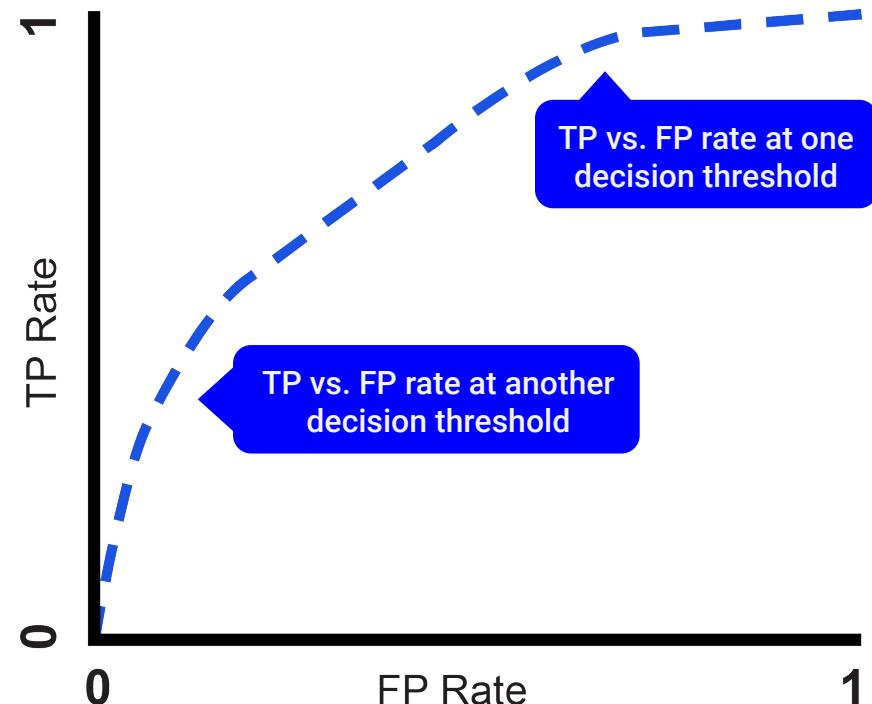
The **FPR** is calculated as follows:

$$FPR = \frac{FP}{FP + TN}$$

Introducing the ROC Curve and AUC

Every point in the ROC curve represents the TPR vs. FPR at different thresholds.

This image shows a typical ROC curve.



Understanding AUC

AUC

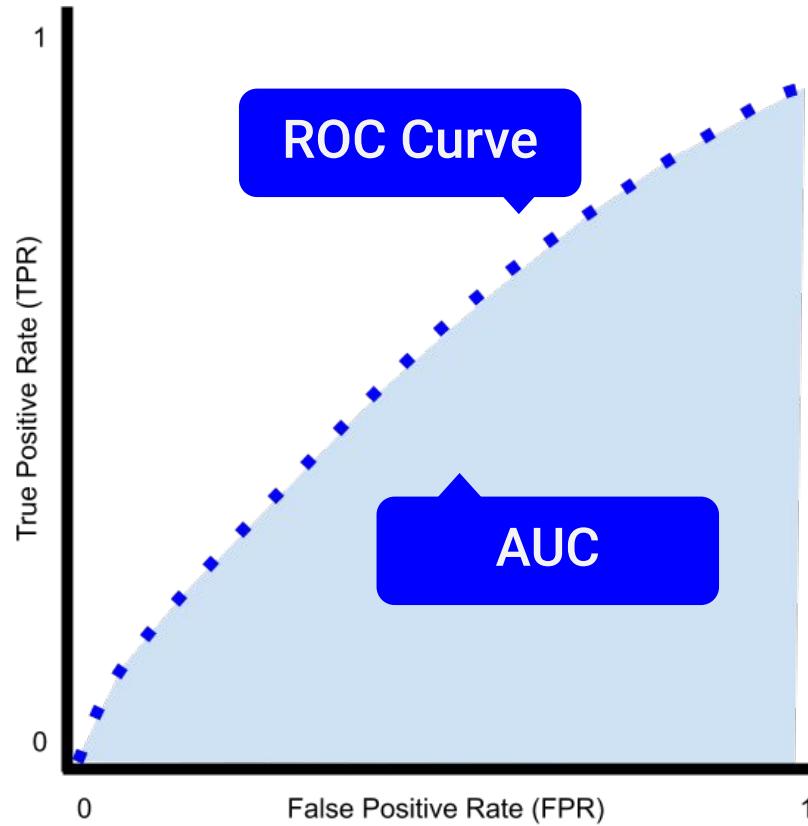
stands for

Area Under the ROC Curve

Introducing the ROC Curve and AUC

Interpreting the ROC curve can be challenging.

Fortunately, we have the AUC that measures the area below the entire ROC curve from (0,0) to (1,1).

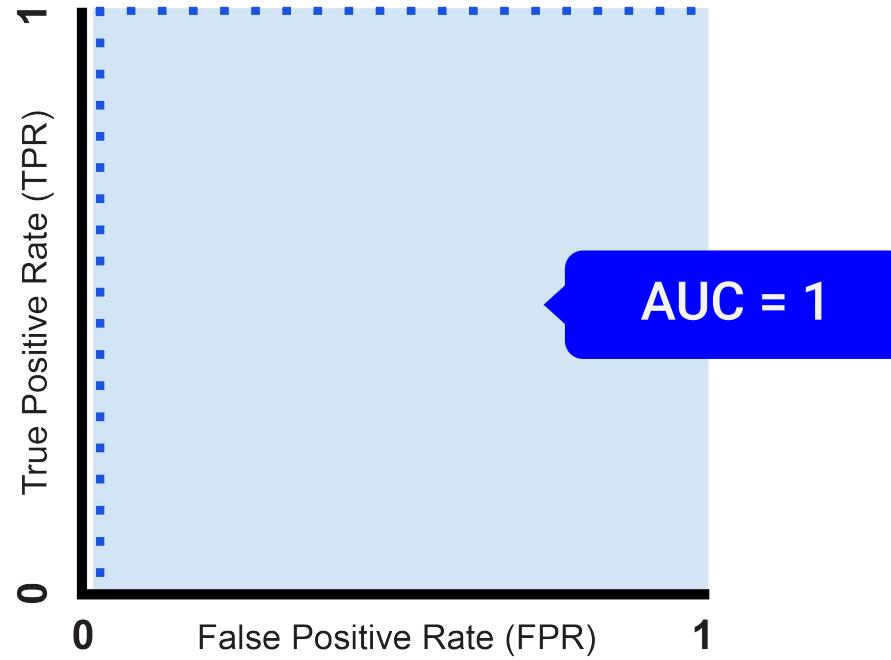




The value of AUC
ranges from 0 to 1.

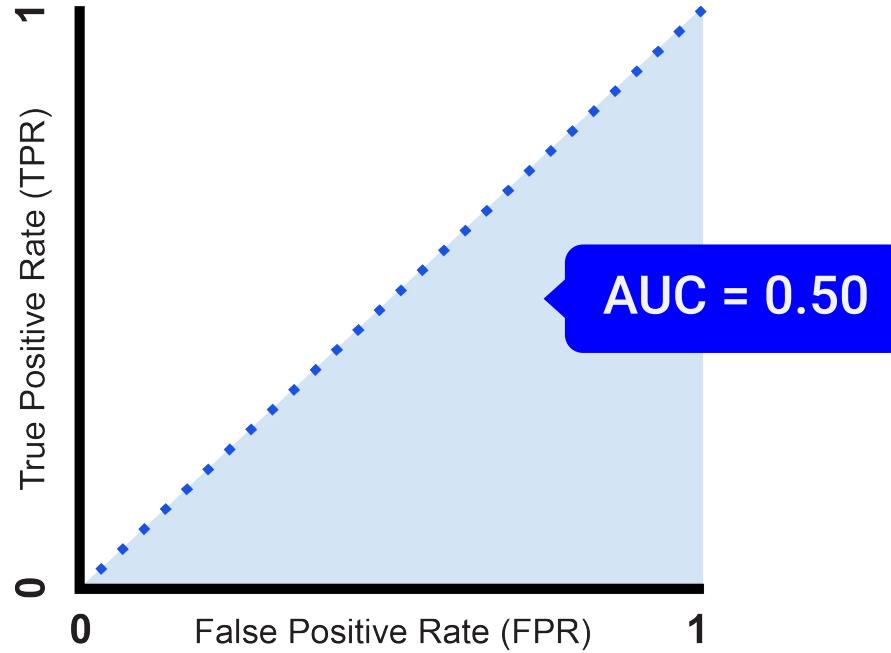
Introducing the ROC Curve and AUC

AUC = 1 is a paradox: it indicates that the model is perfect, but you may not trust your model because it might be overfitted.



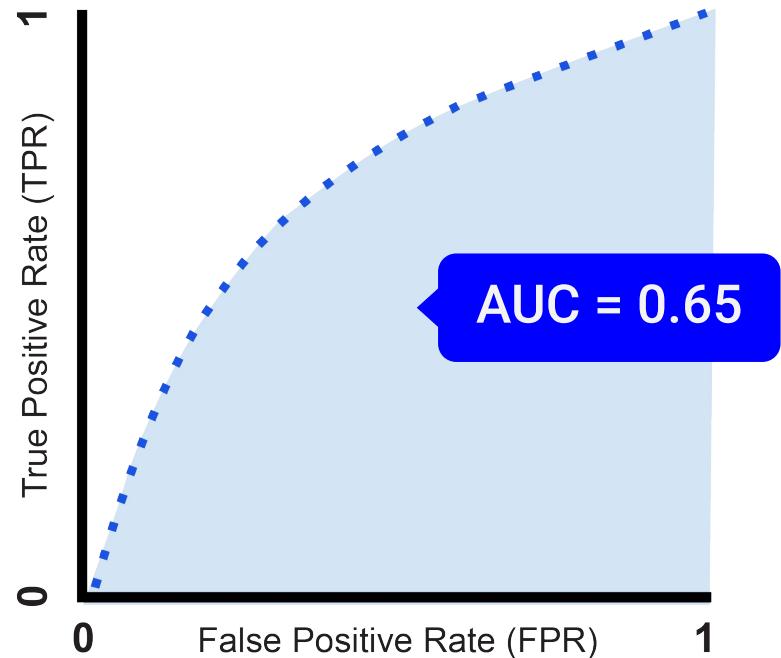
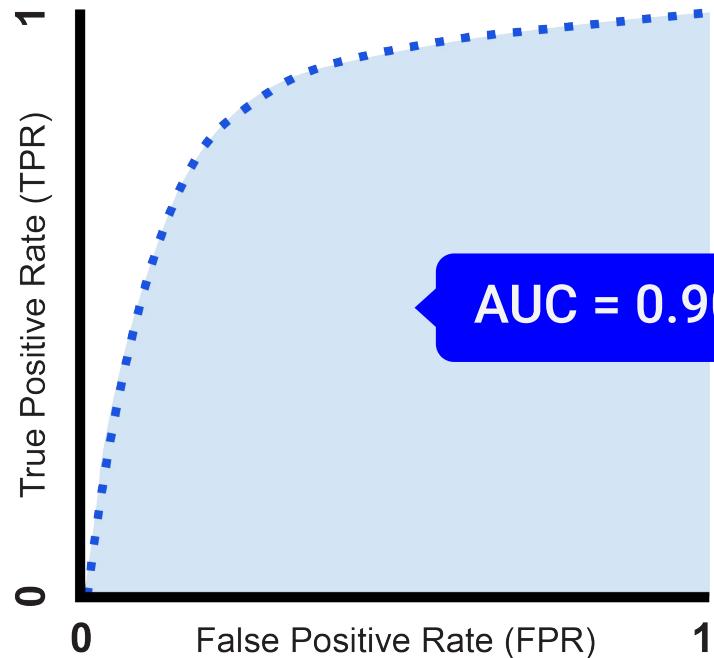
Introducing the ROC Curve and AUC

AUC = 0.50 means that the model is unable to distinguish between positive and negative classes.



Introducing the ROC Curve and AUC

Ideally, we want AUC values ranging between 0 and 1. The higher the AUC, the better the model is at predicting 0s as 0s and 1s as 1s.





So, a model with an
AUC=0.90 may be better
than a model with an
AUC=0.65.

Intro to Recurrent Neural Networks (RNNs)

Intro to Recurrent Neural Networks (RNNs)

RNNs are able to remember the past, and their decisions are influenced by what they have learned in the past.





What Are Recurrent Neural Networks Used For?

What Are Recurrent Neural Networks Used For?

RNNs are used for the following:

01

Natural language processing (NLP)

02

DNA sequences

03

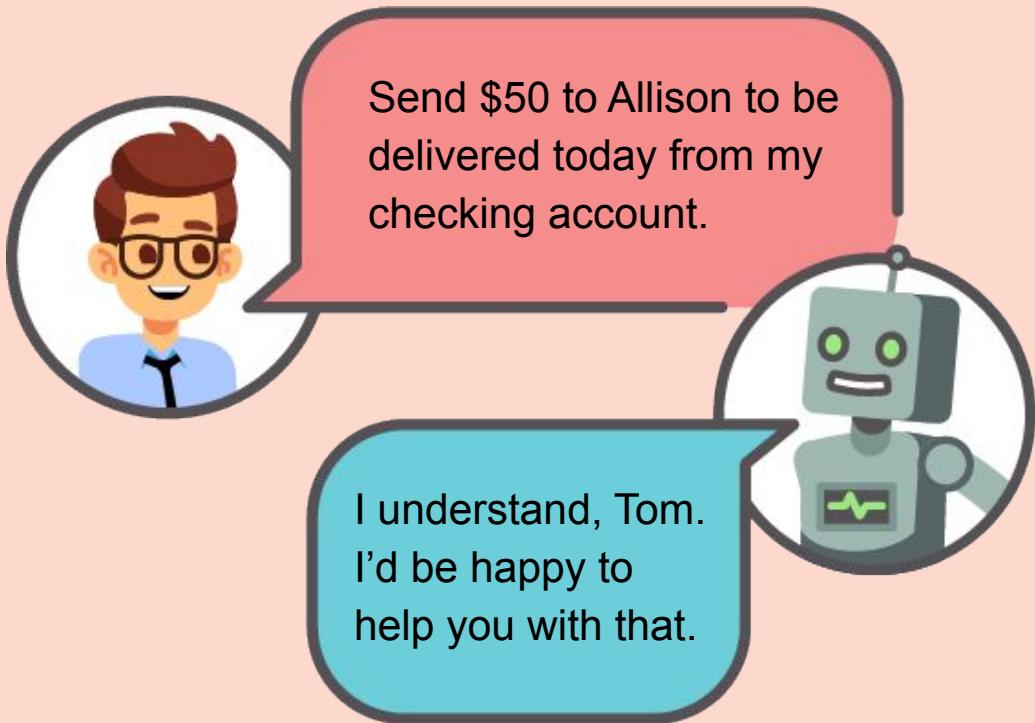
Time series data

04

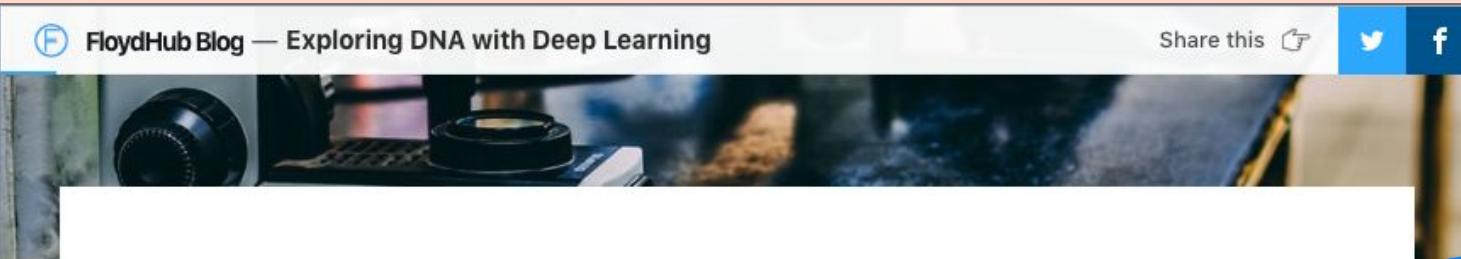
Music composition

What Are Recurrent Neural Networks Used For?

Natural
Language
Processing
(NLP)



What Are Recurrent Neural Networks Used For?



FloydHub Blog — Exploring DNA with Deep Learning

Share this

[Twitter](#) [Facebook](#)

Neural networks are changing the way that Lex Flagel studies DNA.

Lex's recent paper – *The Unreasonable Effectiveness of Convolutional Neural Networks in Population Genetic Inference* – demonstrates how simple deep learning techniques can be used to tackle the ever-changing field of DNA research.

DNA sequences

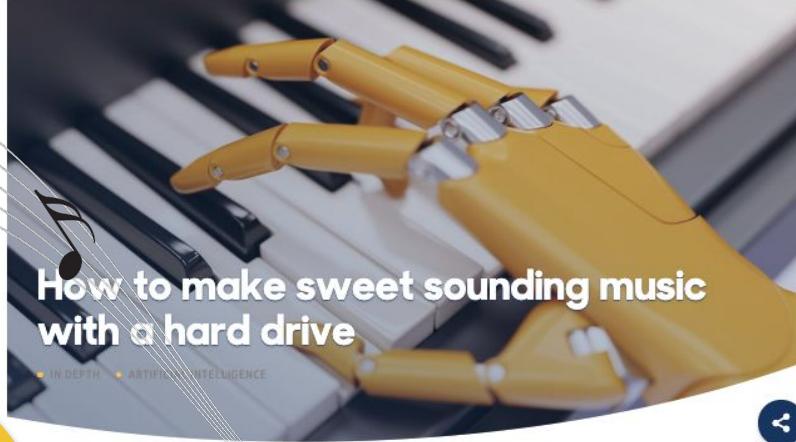
What Are Recurrent Neural Networks Used For?



BBC News Sport Reel Worklife Travel Future More Search 

FUTURE What is BBC Future? Latest Best of... Made on Earth Japan 2020 More 

YOU'RE READING



How to make sweet sounding music with a hard drive

IN DEPTH ARTIFICIAL INTELLIGENCE

Artificial intelligence is being used by musicians to help compose melodies, write lyrics and even perform. It may only be a matter of time before a computer has a number one hit.

By Sarah Griffiths
17th December 2018

A sked how The Beatles approached songwriting, John Lennon quipped "on the M1 (motorway) – turn right, past London." His songwriting partner, Paul McCartney described the process as more of a long and winding road, in which the pair looked for chord shapes and then worked out a melody as if they were "doing a crossword puzzle".

What Are Recurrent Neural Networks Used For?

The screenshot shows a blog post from DeepMind's website. The header includes the DeepMind logo, a navigation bar with links to About, Research, Impact, Blog, Safety & Ethics, and Careers, and a search icon. The main content area has a pink sidebar containing a 'BLOG POST RESEARCH' badge, the date '20 JUL 2016', and the title 'DeepMind AI Reduces Google Data Centre Cooling Bill by 40%'. Below the title are 'SHARE' and 'AUTHORS' sections with social media icons and a blue circular profile. The main body of the post features a large photograph of a data center's cooling infrastructure with many colorful pipes and machinery. A large pink circle on the right side contains the text 'Time Series Data'.

DeepMind

Blog

DeepMind AI Reduces Google Data Centre Cooling Bill by 40%

BLOG POST
RESEARCH

20 JUL 2016

DeepMind AI Reduces Google Data Centre Cooling Bill by 40%

SHARE

AUTHORS

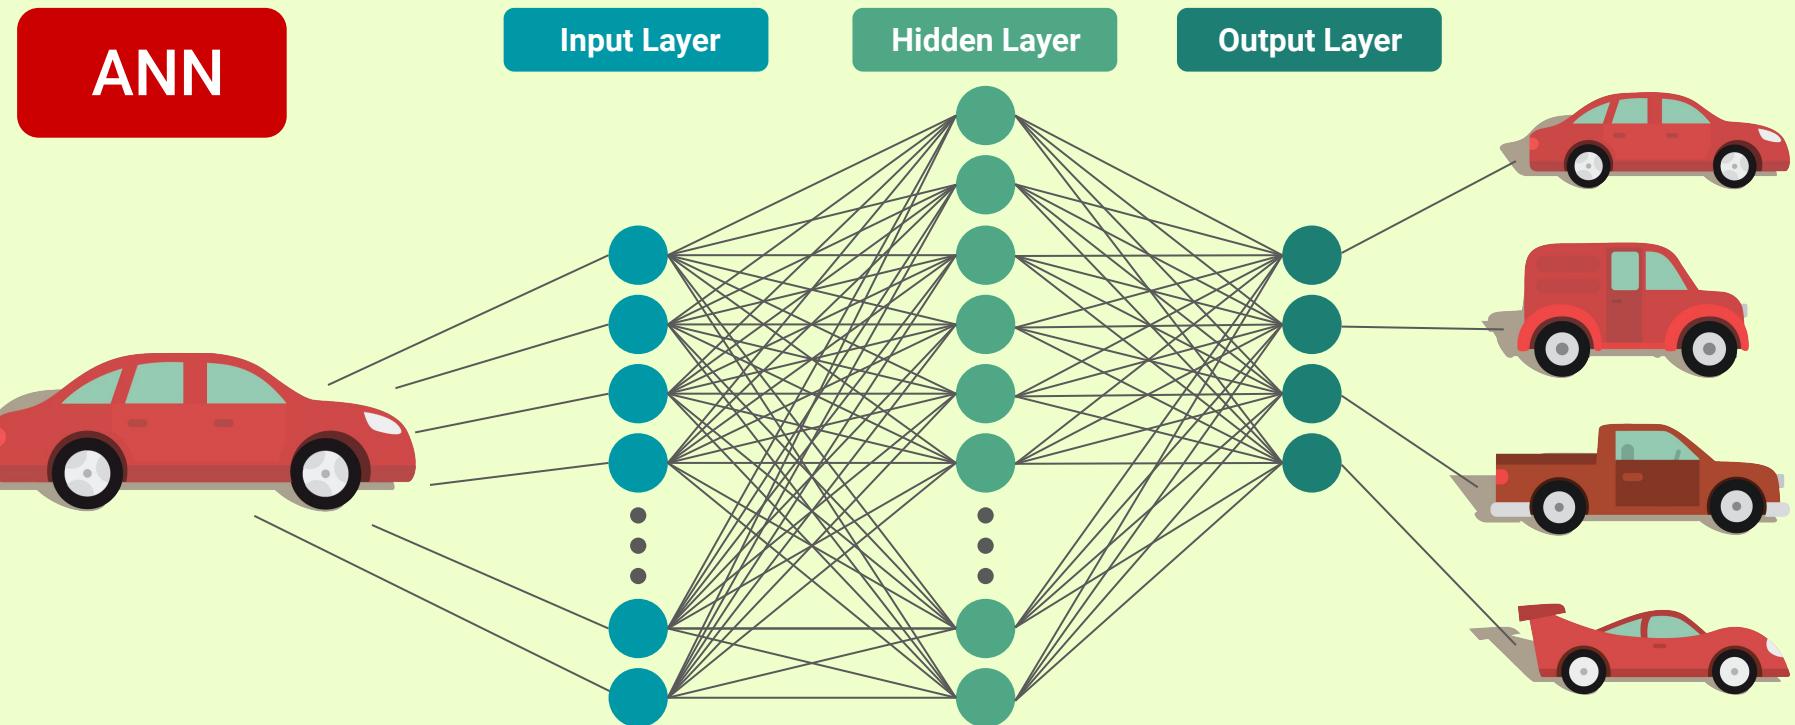
From smartphone assistants to image recognition and translation, machine learning already helps us in our everyday lives. But it can also help us to tackle some of the world's most challenging physical problems – such as energy consumption. Large-scale

Time Series Data

ANNs vs. RNNs

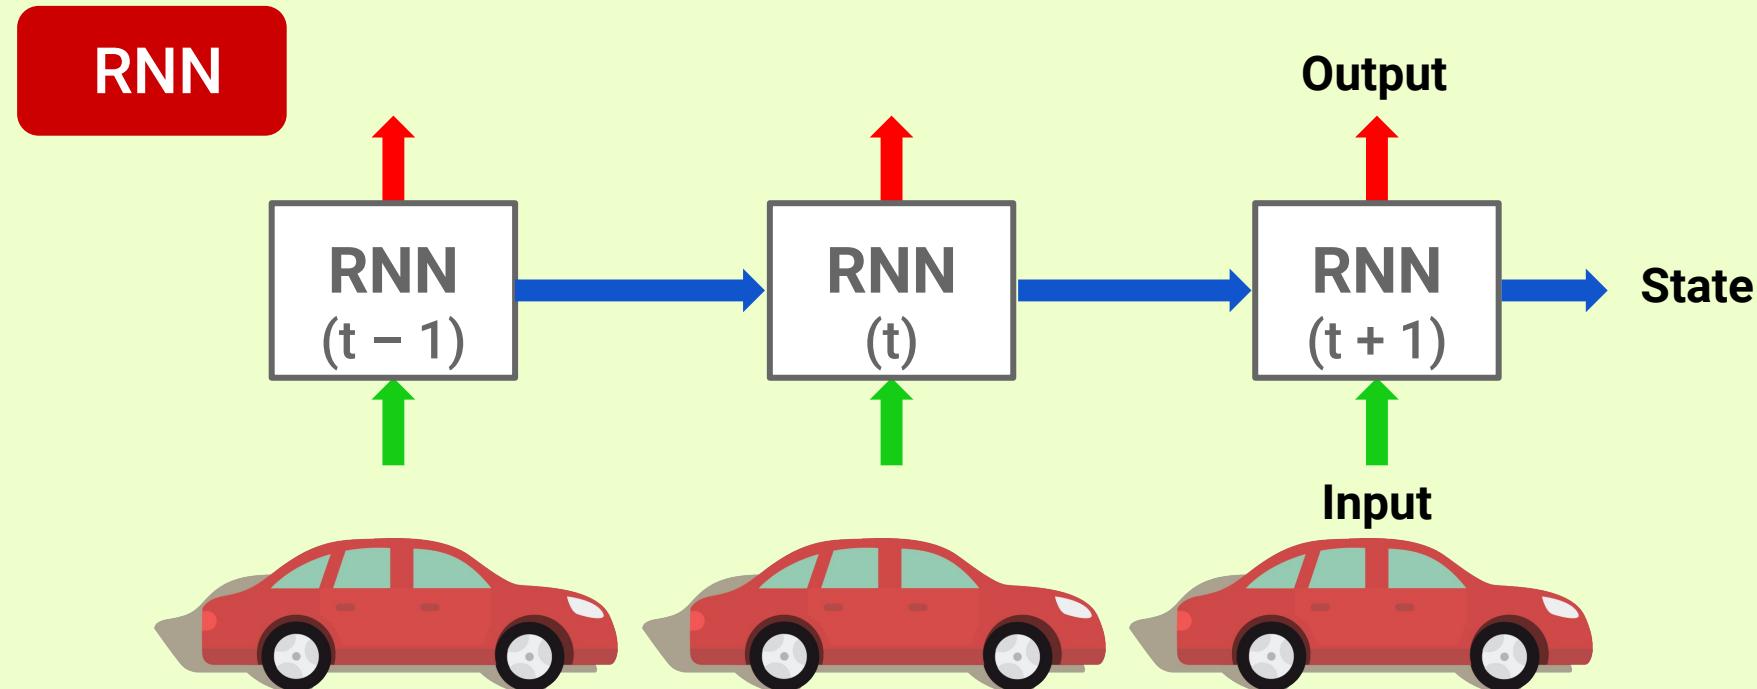
ANNs vs. RNNs

We can use **ANNs** to identify the type of car from a still image.
But can we predict the direction of a car in movement?



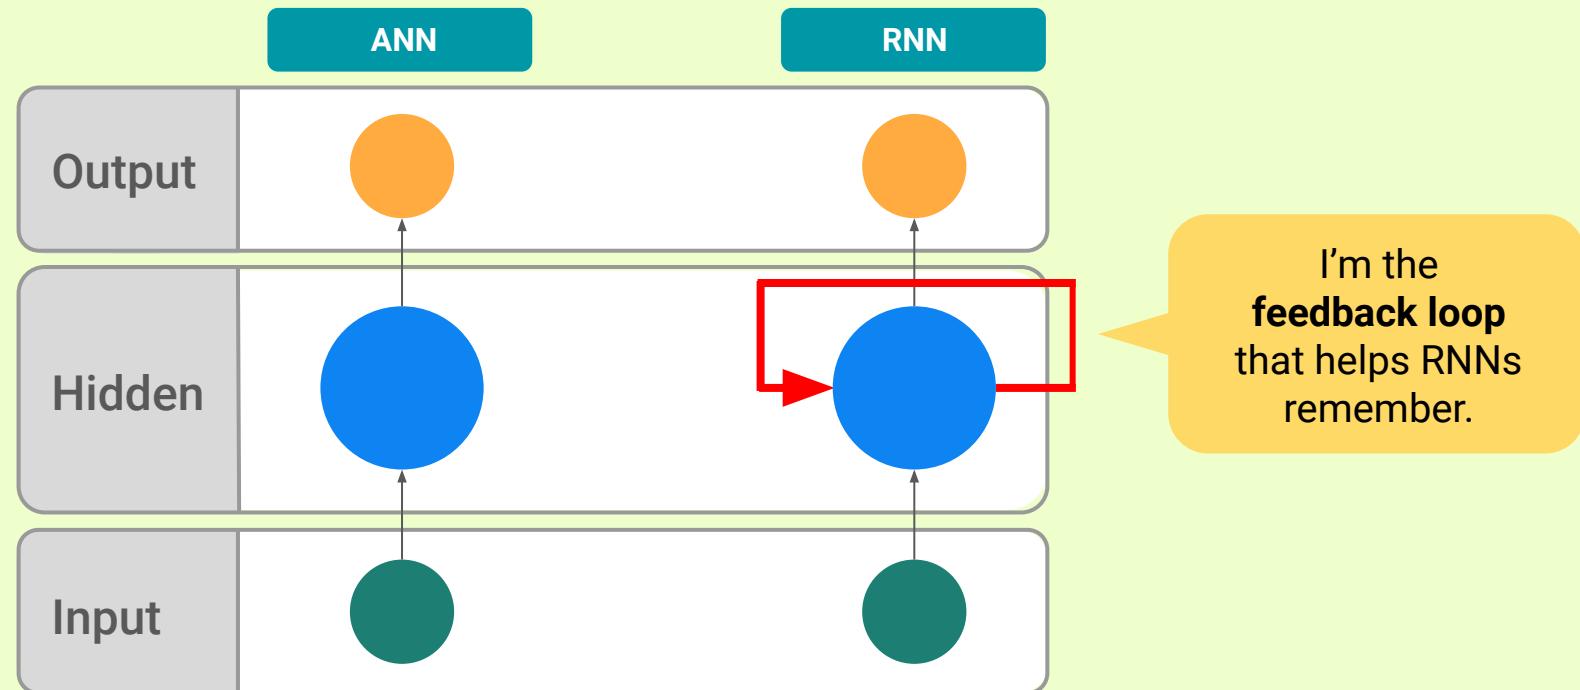
ANNs vs. RNNs

RNNs are good at modeling sequence data thanks to their **sequential memory**. Using RNNs, we can predict that the car is moving to the right.



ANNs vs. RNNs

RNNs are good at modeling sequence data thanks to their **sequential memory**. Using RNNs, we can predict that the car is moving to the right.

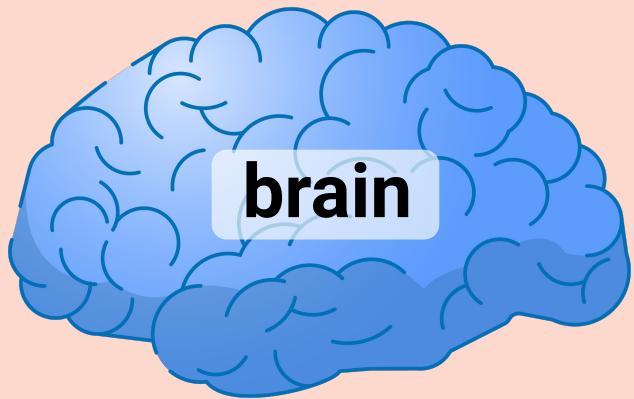




How Do RNNs Work?

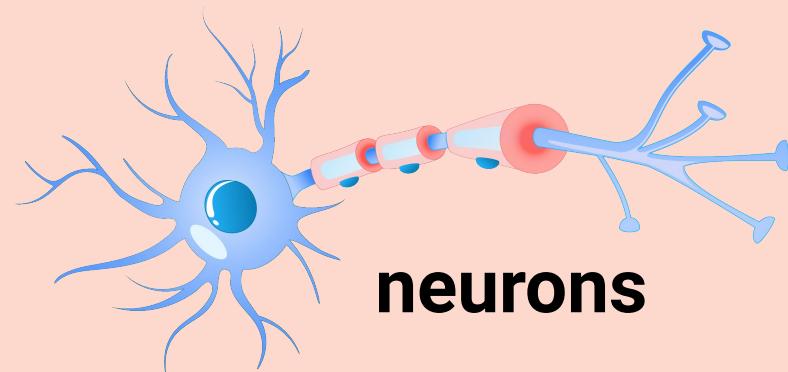
How Do RNNs Work?

When you read this sentence, your



is able to decode it and understand . . .

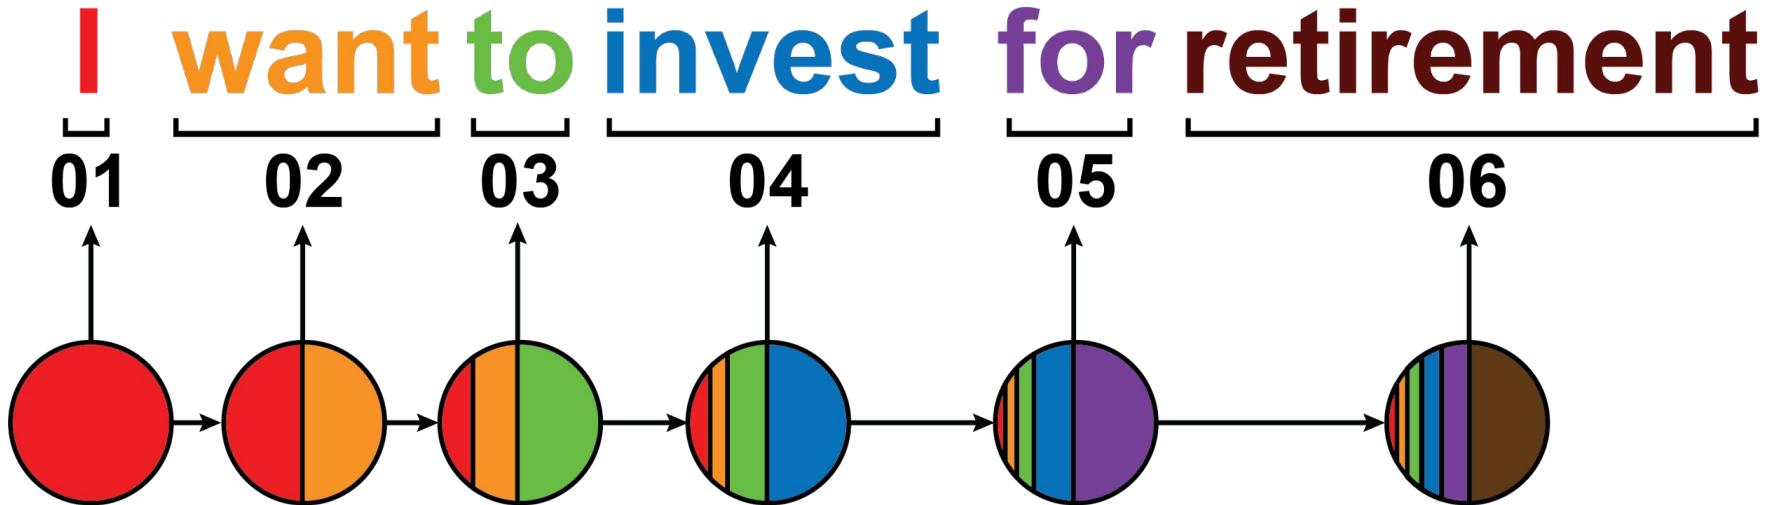
. . . because our



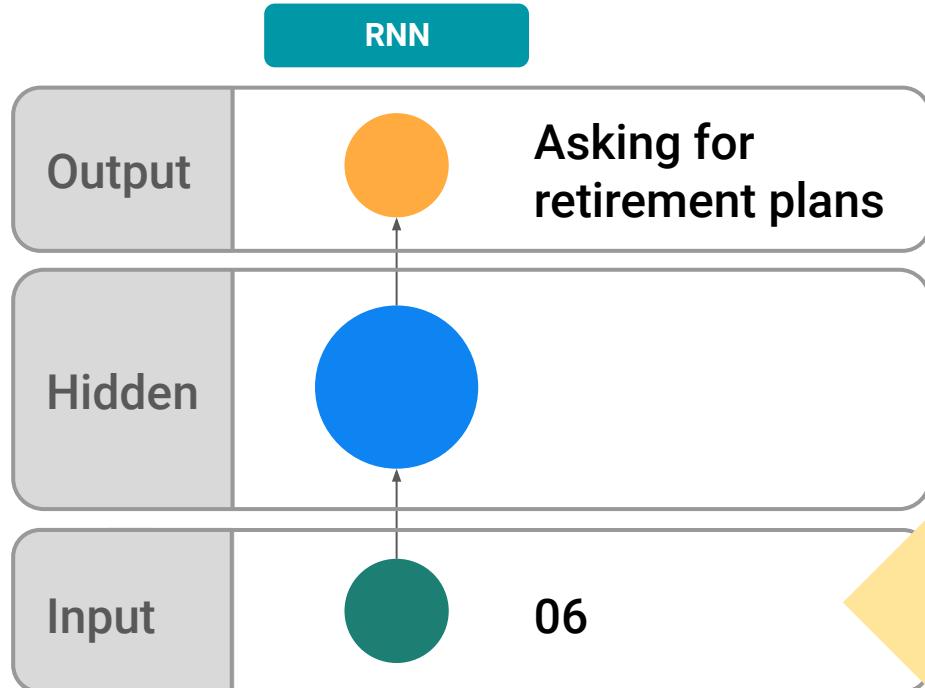
have memory like RNNs.

How Do RNNs Work?

The sentence is split into individual words. RNNs work sequentially, so we feed it one word at a time. By the final step, the RNN has encoded information from all the words in previous steps.



How Do RNNs Work?

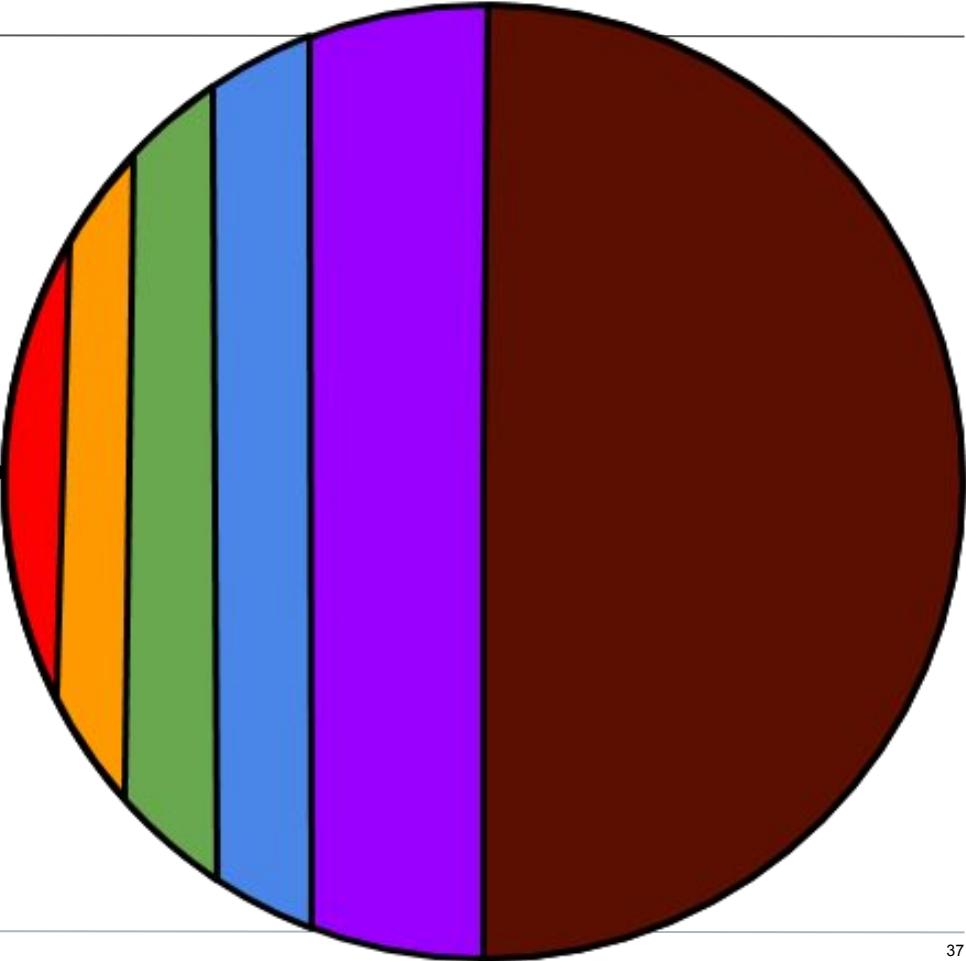


The final output was created from the rest of the sequence. To predict what the phrase means, we take the input and pass it to the feed-forward layer of the RNN to classify the intent.

RNNs Are Forgetful

RNNs only “remember”
the most recent few steps.

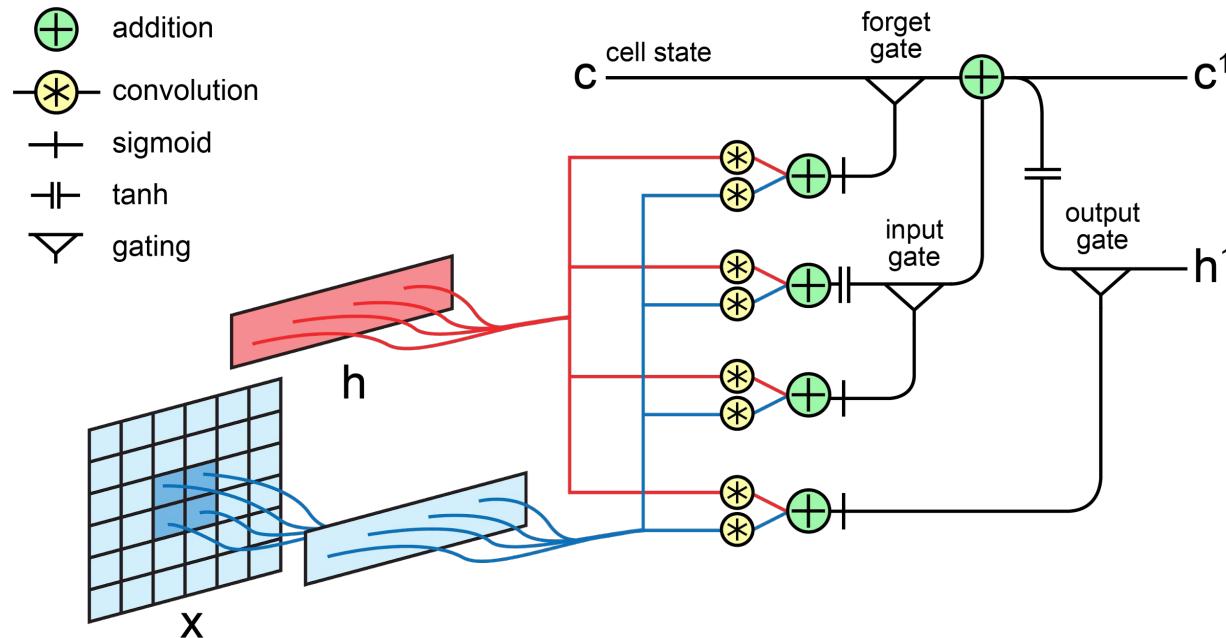
The vanishing gradient in the hidden
states illustrates an issue with
RNNs: **short-term memory**.

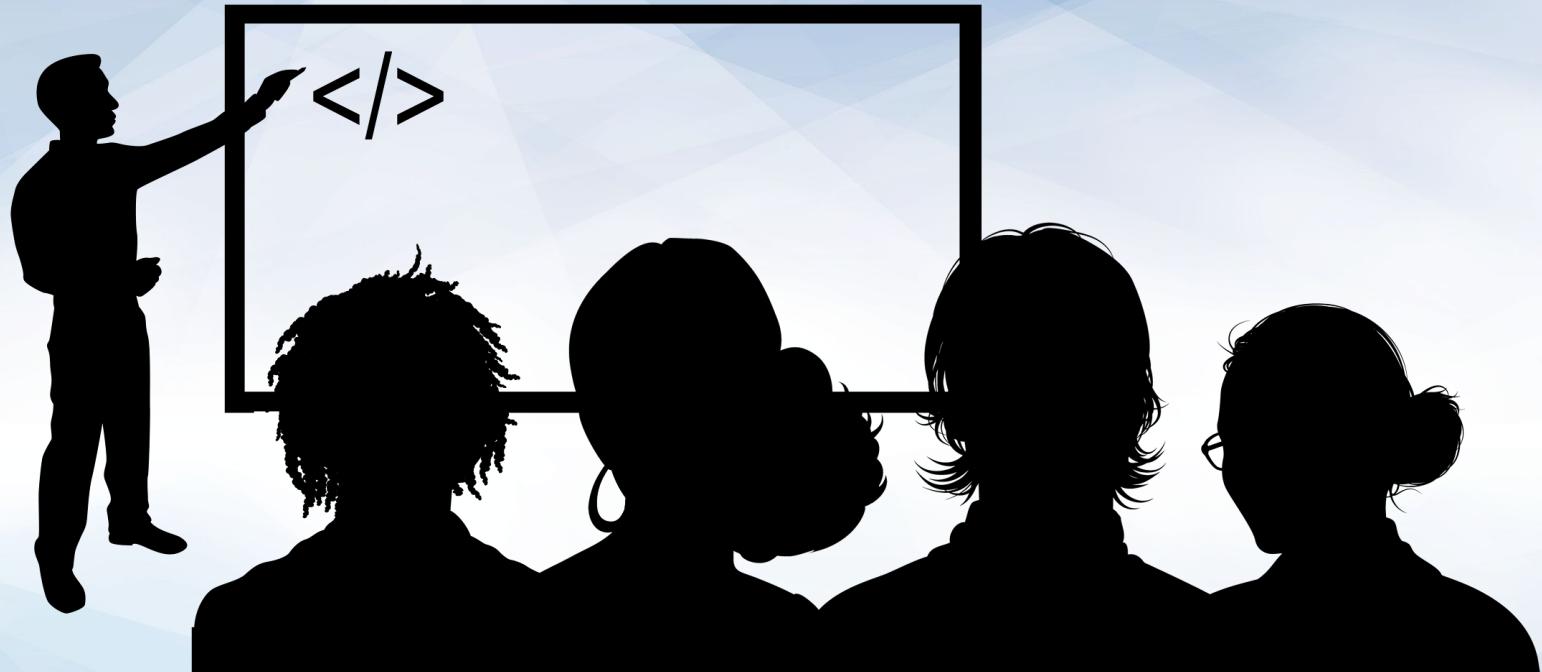


Long Short-Term Memory (LSTM)

LSTMs to the Rescue

LSTM (Long Short-Term Memory) RNNs are one solution for longer time windows. An LSTM RNN works like an original RNN, but it selects which types of longer-term events are worth remembering and which are okay to forget.





Instructor Demonstration
Automatic Text Generation with RNNs

Automatic Text Generation with RNNs

In this demo, we will explore how an RNN can be used to automatically generate text.

Talk to Transformer

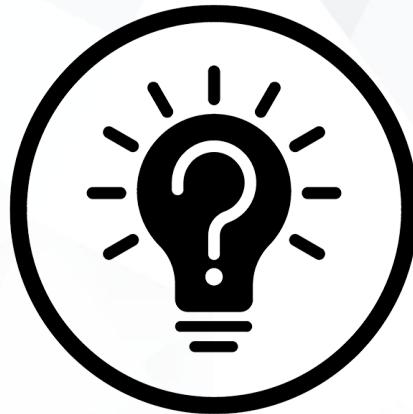
See how a modern neural network completes your text. Type a custom snippet or try one of the examples. [Learn more](#) below.

 Follow @AdamDanielKing for more neat neural networks.

Custom prompt ▾

Type something and a neural network will guess what comes next.

COMPLETE TEXT



Want to Learn More About RNNs?

Take a Look at This Recurrent Neural Networks Cheat Sheet

Shervine Amidi About Projects Teaching Blog About Afshine Amidi MIT

Recurrent Neural Networks

Overview

- Architecture structure
- Applications of RNNs
- Loss function
- Backpropagation

Handling long term dependencies

- Common activation functions
- Vanishing/exploding gradient
- Gradient clipping
- GRU/LSTM
- Types of gates
- Bidirectional RNN
- Deep RNN

Learning word representation

- Notations
- Embedding matrix
- Word2vec
- Skip-gram
- Negative sampling
- GloVe

Comparing words

- Cosine similarity
- t-SNE

Would you like to see this cheatsheet in your native language? You can help us [translating it](#) on GitHub!

CS 230 - Deep Learning English فارسی Français 日本語 한국어 Türkçe

Convolutional Neural Networks Recurrent Neural Networks Tips and tricks

Recurrent Neural Networks cheatsheet

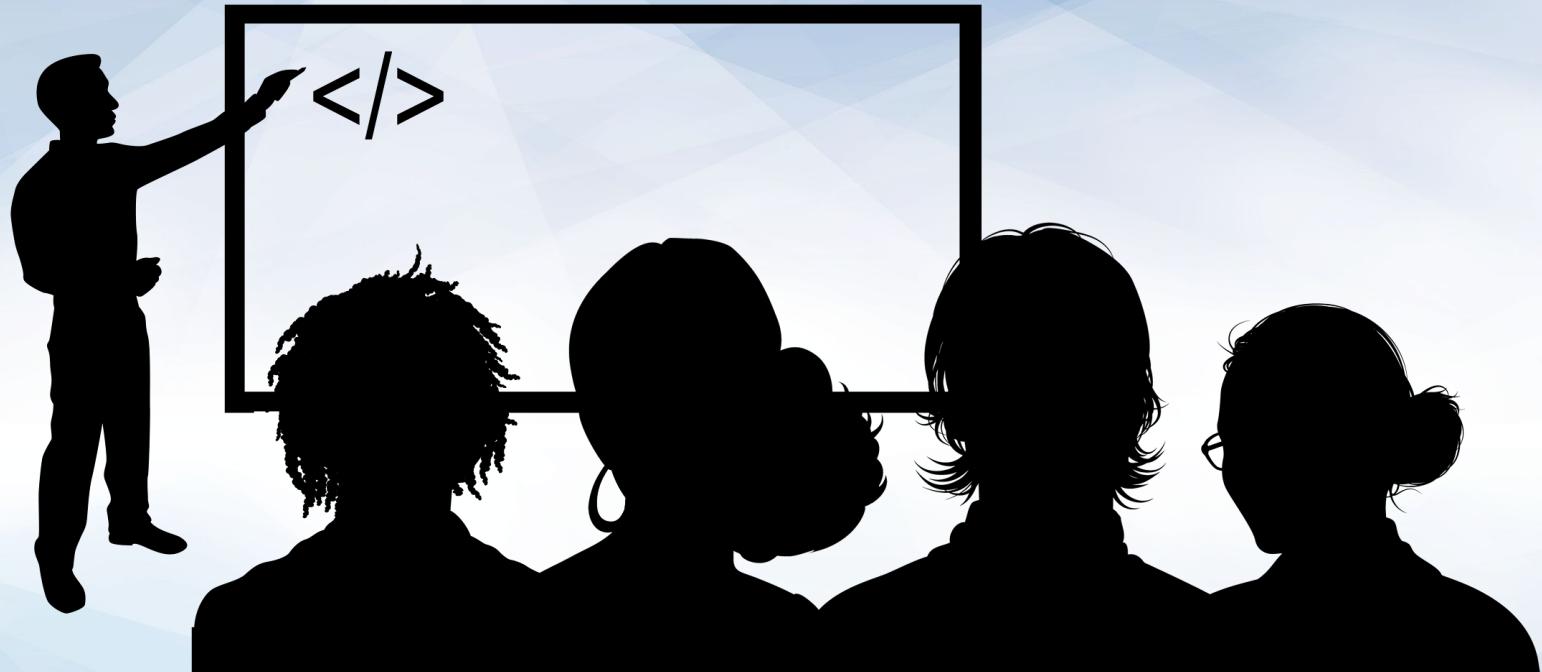
By Afshine Amidi and Shervine Amidi ★ Star 3,793

Overview

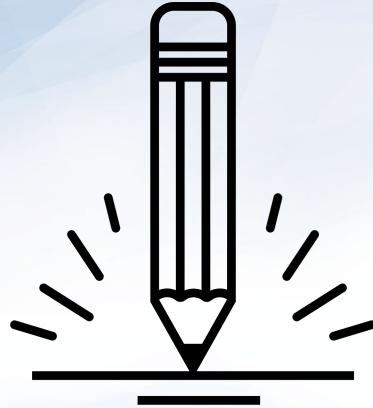
Architecture of a traditional RNN — Recurrent neural networks, also known as RNNs, are a class of neural networks that allow previous outputs to be used as inputs while having hidden states. They are typically as follows:

For each timestep t , the activation $a^{<t>}$ and the output $y^{<t>}$ are expressed as follows:

$$a^{<t>} = g_1(W_{aa}a^{<t-1>} + W_{ax}x^{<t>} + b_a) \quad \text{and} \quad y^{<t>} = g_2(W_{ya}a^{<t>} + b_y)$$



Instructor Demonstration RNNs for NLP - Sentiment Analysis



Activity: RNNs Vs. Vader

In this activity, you will use two different models to score sentiment.

Suggested Time:
40 minutes





Time's Up! Let's Review.

Break





Instructor Demonstration RNN LSTM and Time Series

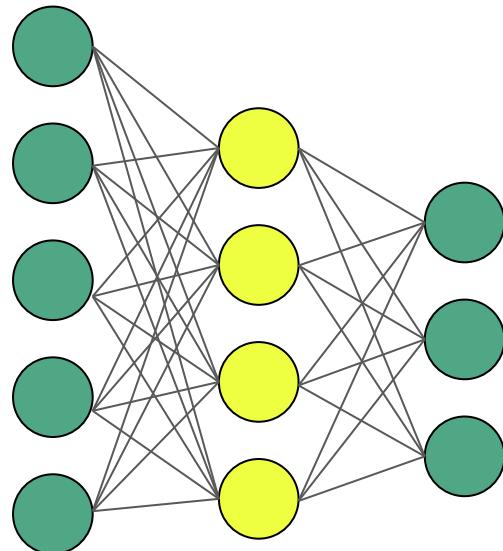
Introduction to Dropout

Introduction to Dropout

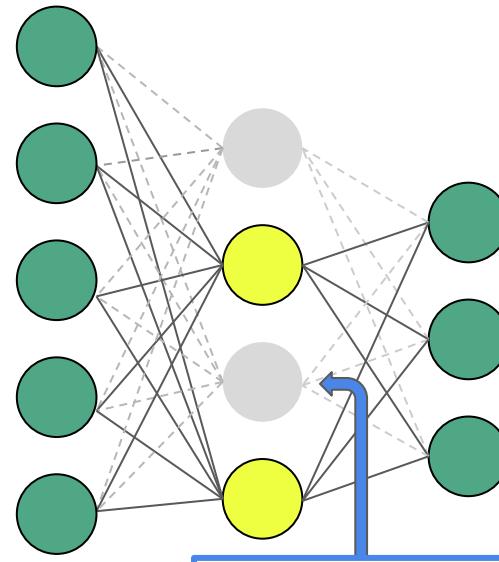
Dropout consists of removing units from the hidden layers by randomly selecting a fraction of the hidden nodes and setting their output to zero, regardless of the input.

A different subset of units is randomly selected every time we feed a training example.

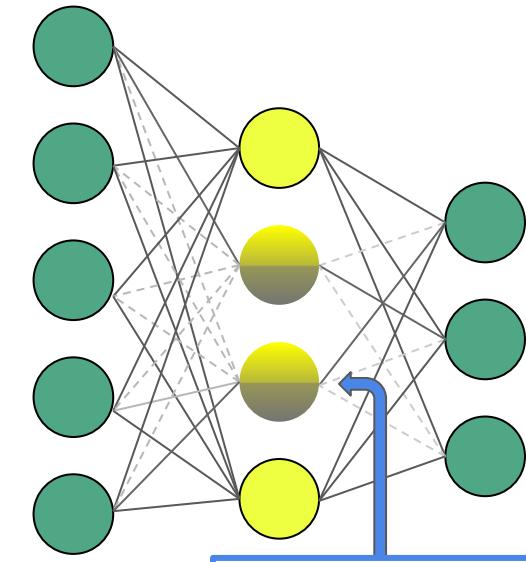
Regular RNN

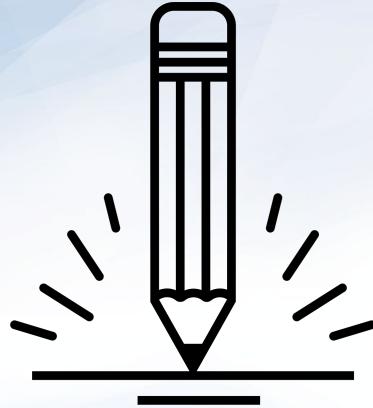


Dropout



DropConnect





Activity: Predicting Gold Closing Prices

In this activity you will gain hands-on experience building an RNN LSTM model for time-series data predicting gold closing prices.

Suggested Time:
30 minutes





Time's Up! Let's Review.

Questions?

*The
End*